

Spring 2024 CSE 380

1-23-2024

Office Hours Update

- Abigail's office hours
 - Mondays: 10:30-12:00pm (3203 EB)
 - Wednesdays: 10:30-12:00pm (3203 EB)
 - Fridays: 10:30-11:30am (always Zoom)
 - Abigail's Zoom link is on Piazza
- Professor Mariani's office hours (always 3145 EB)
 - Mondays: 1:30pm-3:00pm
 - Wednesdays: 2:00pm-3:30pm

Quiz Reminder

- We have our first quiz on Thursday (January 25rd)
- Please remember, no phones, cheatsheets, computers, etc. of any kind
- Make sure you're studying and practicing
- Quiz 2 might be moved, will keep you updated

Project Updates

- Project 1 Released on Feb 16, Due March 6
- Project 2 Released March 8, Due March 20
- Project 3 Released March 22, Due April 3
- Project 4 Released April 5, Due April 17
- All projects will be done in Flask (Python)
 - I will do a small Flask review, and a review of SQLite in Python before Project 1
 - If you want to get a head start, can look up some tutorials on Flask

Review

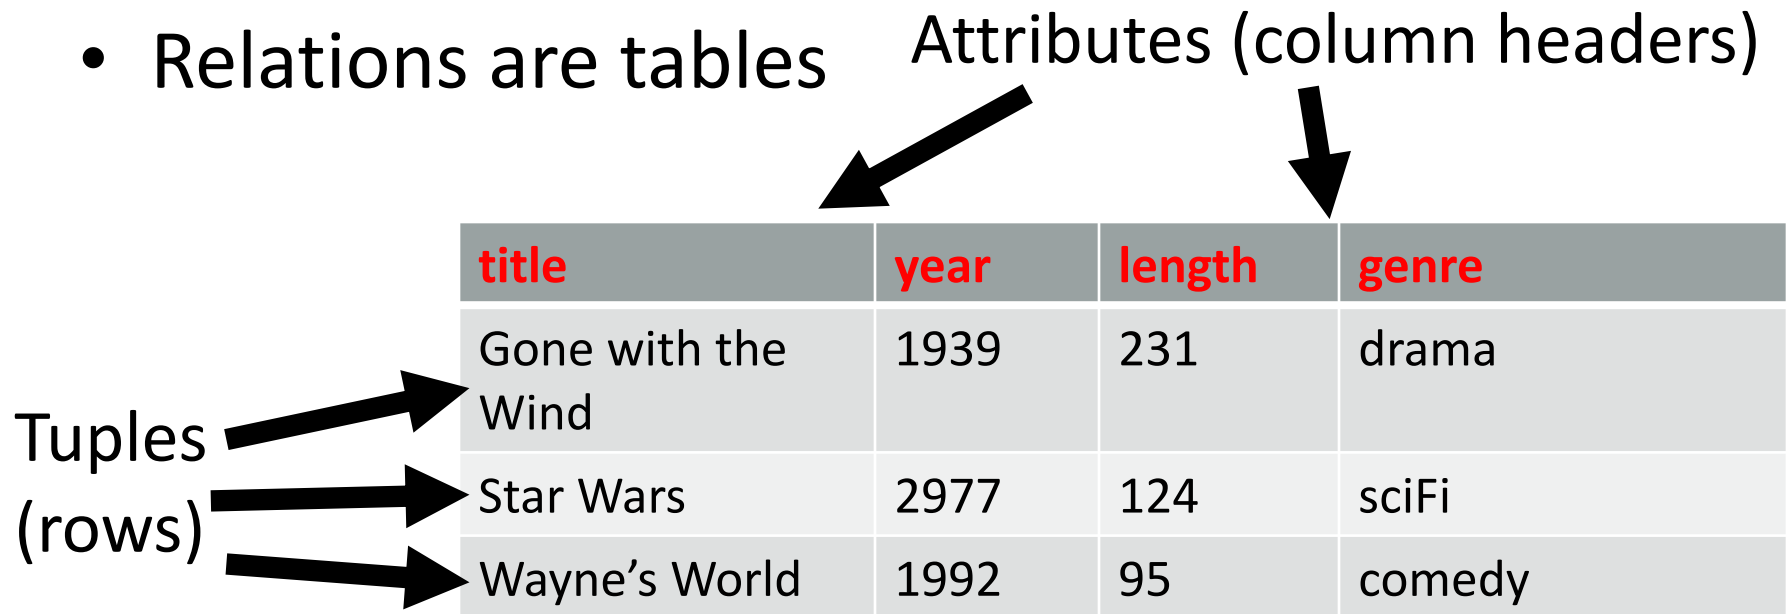
Introduction to SQL

- Data Models – Components
 - Mathematical representation of the data
 - Relational Models (CSV) = tables
 - Semi-structured Models (XML, JSON) = trees/graphs
 - Operations on data (what is allowed, i.e. comparisons, equality, math)
 - Constraints (what types of data are allowed, and where)

Introduction to SQL

- Relations are tables

Attributes (column headers)



title	year	length	genre
Gone with the Wind	1939	231	drama
Star Wars	2977	124	sciFi
Wayne's World	1992	95	comedy

Tuples (rows)

Relation name → Movies

Introduction to SQL

- Schemas (sound familiar?)
 - Relation schema = relation name and attribute list
 - Optionally: types of attributes
 - Example: Movies (title, year, length genre) or Movies (title TEXT, year INTEGER, length INTEGER, genre TEXT)
 - Database = collection of relations
 - Database schema = set of all relation schemas in the database

Introduction to SQL

- Why Relations (tables)?
 - Very simple model
 - Often (but not always) matches how we think about data
 - Abstract model that underlies SQL, the most important database language today

Basic SQL Querying

- Example: Create Table

```
CREATE TABLE Movies (  
    title TEXT,  
    year INTEGER,  
    length INTEGER,  
    genre TEXT  
);
```

Basic SQL Querying

- SQL Values
 - Integers and real values (floats) are represented as you would expect
 - Strings are as well, except they require single quotes
 - Two single quotes = one real quote
 - 'Joe''s Bar' = Joe's Bar
 - Any value can be NULL
 - For now, we will learn later how this can be prevented

Basic SQL Querying

- Insertion

- To insert a single record:

- ```
INSERT INTO <relation>
```

- ```
VALUES ( <list of values> );
```

- Example: add Star Wars to our Movies relation

- ```
INSERT INTO Movies
```

- ```
VALUES ('Star Wars', 1977, 124, 'sciFi');
```

Basic SQL Querying

- Select
 - Using our Movies relation, what movies have been released?

SELECT title

FROM Movies;

Column we want to
select





Table we want to
select from



Basic SQL Querying

- Select and Order By
 - Using our Movies relation, what movies came out, ordered from oldest to newest

SELECT title

FROM Movies

ORDER BY year;

Column we want to
select





Table we want to
select from



Column to order
results by



Basic SQL Querying

- Select and Order By
 - Using our Movies relation, what movies were released, ordered by release date

```
SELECT title  
FROM Movies  
ORDER BY year, length;
```

If same year, break
ties by 'length'
column



Basic SQL Querying

- Select Multiple Attributes
 - Using our Movies relation, what are the names of the movies and the year they were released

```
SELECT title, year  
FROM Movies  
ORDER BY year;
```


Basic SQL Querying

- Select *
 - What are all of the attributions of the relation 'players'?

```
SELECT *
```


```
FROM Movies
```

```
ORDER BY year;
```

SQLite Demo

- Download sqlite
- Type:
 - `sqlite3 test.db`

If you're starting a new DB, you can name this whatever you want



```
host-262820:~ cse498$ sqlite3 test.db
SQLite version 3.30.0 2019-10-04 15:03:17
Enter ".help" for usage hints.
sqlite> █
```

SQLite Demo

- Create the table

```
host-262820:~ cse498$ sqlite3 test.db
SQLite version 3.30.0 2019-10-04 15:03:17
Enter ".help" for usage hints.
sqlite> CREATE TABLE Movies(title TEXT, year INTEGER, length INTEGER, genre TEXT);
sqlite> █
```

- Insert values into the table

```
host-262820:~ cse498$ sqlite3 test.db
SQLite version 3.30.0 2019-10-04 15:03:17
Enter ".help" for usage hints.
sqlite> CREATE TABLE Movies(title TEXT, year INTEGER, length INTEGER, genre TEXT);
sqlite> INSERT INTO Movies VALUES ('Gone With the Wind', 1939, 231, 'drama');
sqlite> INSERT INTO Movies VALUES ('Star Wars', 1977, 123, 'sciFi');
sqlite> INSERT INTO Movies VALUES ('Wayne's World', 1992, 95, 'comedy');
sqlite> █
```

SQLite Demo

- Let's see what's in the DB

```
[sqlite> SELECT * FROM Movies;  
Gone With the Wind|1939|231|drama  
Star Wars|1977|123|sciFi  
Wayne's World|1992|95|comedy  
sqlite> █
```

- That's Ugly

```
[sqlite> .headers on  
[sqlite> .mode columns
```

SQLite Demo

- One more SELECT

```
sqlite> .headers on
sqlite> .mode columns
sqlite> SELECT * FROM Movies;
title                year          length        genre
-----
Gone With the Wind   1939          231           drama
Star Wars            1977          123           sciFi
Wayne's World        1992          95            comedy
sqlite> █
```

SQLite Demo **NEW**

- Saving the SQL

```
sqlite> .dump
PRAGMA foreign_keys=OFF;
BEGIN TRANSACTION;
CREATE TABLE Movies(title TEXT, year INTEGER, length INTEGER, genre TEXT);
INSERT INTO Movies VALUES('Gone With the Wind',1939,231,'drama');
INSERT INTO Movies VALUES('Star Wars',1977,123,'sciFi');
INSERT INTO Movies VALUES('Wayne''s World',1992,95,'comedy');
COMMIT;
sqlite> .output test.sql
```

New Material

WHERE Clauses

- `SELECT ... FROM ... WHERE ...;`
- The where clause stipulates a test (predicate) that each row must pass to be returned in the select statement.

```
SELECT title, genre FROM Movies WHERE length  
< 120;
```


WHERE Clauses

SELECT title FROM Movies WHERE length < 120;

```
[sqlite> SELECT title FROM Movies WHERE length < 120;
title
-----
Wayne's World
Back to the F
```

WHERE Clauses

SELECT title, genre FROM Movies WHERE length
< 120;

```
[sqlite> SELECT title, genre FROM Movies WHERE length < 120;
title          genre
-----
Wayne's World comedy
Back to the F  adventure
```

SQL Operators

- Arithmetic Operators:
 - + - * / %
- Comparison Operators:
 - = != <> > < >= <=
- Logical Operators
 - ALL AND ANY BETWEEN EXISTS IN LIKE NOT OR IS

Between

SELECT title, genre FROM Movies WHERE length
BETWEEN 90 AND 130;

note: between is **inclusive of the bounds**

```
sqlite> SELECT title, genre FROM Movies WHERE length BETWEEN 90 AND 130;
title      genre
-----
Star Wars  sciFi
Wayne's Wo comedy
Back to th adventure
```

Between

SELECT title, genre FROM Movies WHERE length
NOT BETWEEN 90 AND 130;

```
sqlite> SELECT title, genre FROM Movies WHERE length NOT BETWEEN 90 AND 130;
title          genre
-----
Gone With the Wind  drama
The Shawshank Rede  drama
The Godfather       drama
The Dark Knight     action
The Matrix          sciFi
```

Between

- Why can't we do?

```
SELECT title, genre FROM Movies WHERE length  
>= 90 AND length <= 130;
```

IS and IS NOT

SELECT title, genre FROM Movies WHERE length
IS NULL;

'IS' operator is only to be used to test for
NULL values

```
sqlite> SELECT title, genre FROM Movies WHERE length is NULL;
title          genre
-----
Good Will Hunting drama
```

AND and OR

SELECT title, genre FROM Movies WHERE length
BETWEEN 90 AND 130 AND genre = 'comedy';

```
sqlite> SELECT title, genre FROM Movies WHERE length BETWEEN 90 AND 130 AND genre = 'comedy';
```

title	genre
Wayne's World	comedy

An Aside: COUNT()

SELECT COUNT(*) FROM Movies WHERE length BETWEEN 90 AND 130 AND genre = 'comedy';

```
sqlite> SELECT COUNT(*) FROM Movies WHERE length BETWEEN 90 AND 130 AND genre = 'comedy';  
COUNT(*)  
-----  
1
```

SELECT COUNT(*) FROM Movies WHERE genre = 'drama';

```
sqlite> SELECT COUNT(*) FROM Movies WHERE genre = 'drama';  
COUNT(*)  
-----  
4
```

LIKE

SELECT title FROM Movies WHERE title LIKE
'The%';

```
sqlite> SELECT title FROM Movies WHERE title LIKE 'The%';  
title  
-----  
The Shawshank Redemption  
The Godfather  
The Dark Knight  
The Matrix
```

LIKE

SELECT title FROM Movies WHERE title LIKE 'The%';

- % is a wildcard that can match 0 or more characters
 - Probably should have been '*', but already being used
- _ (underscore) is a wildcard representing a single character
 - Should have been '.' but again, already taken

LIKE

SELECT title FROM Movies WHERE title LIKE
'The%' AND genre = 'drama';

```
sqlite> SELECT title FROM Movies WHERE title LIKE 'The%' AND genre = 'drama';  
title  
-----  
The Shawshank Redemption  
The Godfather
```

IN

SELECT title, genre FROM Movies WHERE length
IN(90,130);

- IN tests for inclusion in a static, parenthesized list

```
sqlite> SELECT title, genre FROM Movies WHERE length IN (90, 130);  
sqlite> █
```

Question

- What should happen for a query where nothing matches the WHERE clause, but one of the rows has a NULL in the WHERE column
 - It should return 0 rows
 - It should raise a syntax error
 - It should return 1 row (the row with the NULL)
 - It should raise a "no rows" error

Comparing NULLS to Values

- Comparing any value (including NULL itself) with NULL yields UNKNOWN.
- A tuple is in a query answer if and only if the WHERE clause is TRUE (not FALSE or UNKNOWN).

SQL Column Constraints

- Unique:
 - Used to specify a column whose values will always be unique for every row.
 - CREATE TABLE students (msu_id TEXT UNIQUE,
pid INTEGER UNIQUE,
first_name TEXT);

The DBMS will raise an error if the constraint is violated.

Should NULLS be Allowed in UNIQUE columns?

- No
- Only one
- As many as desired
- Depends on the data type

NOT NULL

- Used to designate columns that may not be NULL

```
CREATE TABLE students (msu_id TEXT UNIQUE NOT  
NULL, section INTEGER NOT NULL);
```

- Raises error if a NULL value is inserted.

Primary key

- Each table can have at most one PRIMARY KEY.
 - It is implicitly NOT NULL and UNIQUE
 - It is used to specify a specific column in the table

```
CREATE TABLE students (msu_id TEXT PRIMARY KEY,  
                        grade REAL);
```

Composite PRIMATE KEY

- You can create a primary key out of multiple columns if the combination is unique

```
CREATE TABLE students (first_name TEXT,  
                        last_name TEXT,  
                        PRIMARY KEY (first_name, last_name));
```

Why Use PRIMARY KEYS?

- They indicate values that can be used to easily connect two relations together (joins).
- They indicate to the database good attributes to organize the rows by.
 - This allows for more optimized queries.
- Best practice is for every table to have a primary key
- Will cover very in depth, later in the course

Which Column(s) Would Make the Best Primary Key for a Students Table?

- Social Security Number
- First and Last Name
- Student Number
- MSU Net ID

Qualified Table Names

- The qualified name of a table is <name of database>.<name of relation> (e.g. main.students)
- The qualified name of a column is <qualified relation>.<name of column> (e.g. main.students.msu_id)
- Sometimes it is simpler to preface column names with their relation (we'll see why shortly)

SELECT name FROM students;

- Is exactly the same as:

SELECT students.name FROM students;

Qualified Table Names

- You can also use the pattern anywhere you use a column name:

```
SELECT * FROM students WHERE students.name = 'James';
```


SQL Alias (AS)

- Another helpful piece of syntax is aliases if you want to refer to a column by a different name.

```
SELECT msu_net_id AS id FROM students  
WHERE id = 'mariani4';
```

- You can also use it on tables:

```
SELECT cse380.msu_net_id FROM students AS cse380;
```

```
SELECT cse380.msu_net_id FROM students cse380;
```

– Implicit Alias, not recommended

Joins

- Often, queries need to combine (join) rows from multiple tables.
- Joins specify which rows get merged with rows from a second table.
- As there are many possible methods to match rows together, there are many types of joins.
- You will need to know (memorize) the different joins.
- Annoying, but very useful knowledge.

Joins

professors

id	name
esfahanian	Abdol
mariani	James
dyksen	Wayne
NULL	demo-prof

```
CREATE TABLE professors (id TEXT, name TEXT);
INSERT INTO professors VALUES('esfahanian','Abdol');
INSERT INTO professors VALUES('mariani','James');
INSERT INTO professors VALUES('dyksen','Wayne');
INSERT INTO professors VALUES(NULL, 'demo-student');
```

associates

responder_id	associate_id
mariani	dyksen
mariani	esfahanian
dyksen	mariani
esfahanian	dyksen

```
CREATE TABLE associates (responder_id TEXT, associate_id TEXT);
INSERT INTO associates VALUES('mariani','dyksen');
INSERT INTO associates VALUES('mariani','esfahanian');
INSERT INTO associates VALUES('dyksen','mariani');
INSERT INTO associates VALUES('esfahanian','dyksen');
```

Cross Join

- Combine every row in one table with every row in the other

```
SELECT * FROM professors CROSS JOIN associates;
```

- Implicit Cross Join (equivalent)

```
SELECT * FROM students, associates;
```

- Also called the Cartesian Product (as it results in every possible combination)
- Not often used, because you rarely need all possible combinations.

Inner Join


- INNER JOIN is like the CROSS JOIN, except only rows that match a predicate are allowed. Frequently that predicate includes primary keys to match associated data

```
SELECT * FROM professors INNER JOIN  
associates  
    ON id = responder_id;
```

- You can equivalently do an implicit inner join and use WHERE to filter rows:

```
SELECT * FROM students, associates  
    WHERE id = responder_id;
```

Inner Join



Returned rows are those that match both tables (shaded).

Inner Join

- `SELECT * FROM professors INNER JOIN associates ON id = responder_id;`

```
sqlite> SELECT * FROM professors INNER JOIN associates ON id = responder_id;
```

id	name	responder_id	associate_id
esfahanian	Abdol	esfahanian	dyksen
mariani	James	mariani	dyksen
mariani	James	mariani	esfahanian
dyksen	Wayne	dyksen	mariani


Full Outer Join

- Like an INNER JOIN, the FULL OUTER JOIN includes all rows that match a predicate. In addition however, rows that don't have a match are included with NULL values where their corresponding data would be

```
SELECT * FROM professors FULL OUTER JOIN associates  
ON id = responder_id;
```

- SQLite doesn't implement FULL OUTER JOIN,): . But I'll show you how to do it anyways in a later lecture.


Full Outer Join



Left Outer Join

- LEFT OUTER JOIN includes all the rows to the left of the JOIN keyword, where there is a match, it add the matching column data to rows. If there isn't a match NULL values are used instead.
- `SELECT * FROM professors LEFT OUTER JOIN associates
ON id = responder_id;`
- Note: order matters!!!
- `SELECT * FROM associates LEFT OUTER JOIN professors
ON id = responder_id;`

Left Outer Join



Left Outer Join

SELECT * FROM professors LEFT OUTER JOIN associates
ON id = responder_id;


```
[sqlite> SELECT * FROM professors LEFT OUTER JOIN associates ON id=responder_id;
id            name            responder_id  associate_id
-----
esfahanian    Abdol            esfahanian    dyksen
mariani       James           mariani       dyksen
mariani       James           mariani       esfahanian
dyksen        Wayne           dyksen        mariani
demo-stude
```

Left Outer Join Without Matches

- If we want to find the rows in the left table that don't have matches, we can do a LEFT OUTER JOIN and filter out the rows with matches.

```
SELECT * FROM professors LEFT OUTER JOIN associates  
      ON id = responder_id WHERE responder_id IS NULL;
```

Left Outer Join Without Matches




What rows should a left inner join return?

- Rows that match the predicate
- Row that match plus the unmatched rows on the left
- Rows that are unmatched and the matches on the left
- ???

Practice

Chinook Dataset

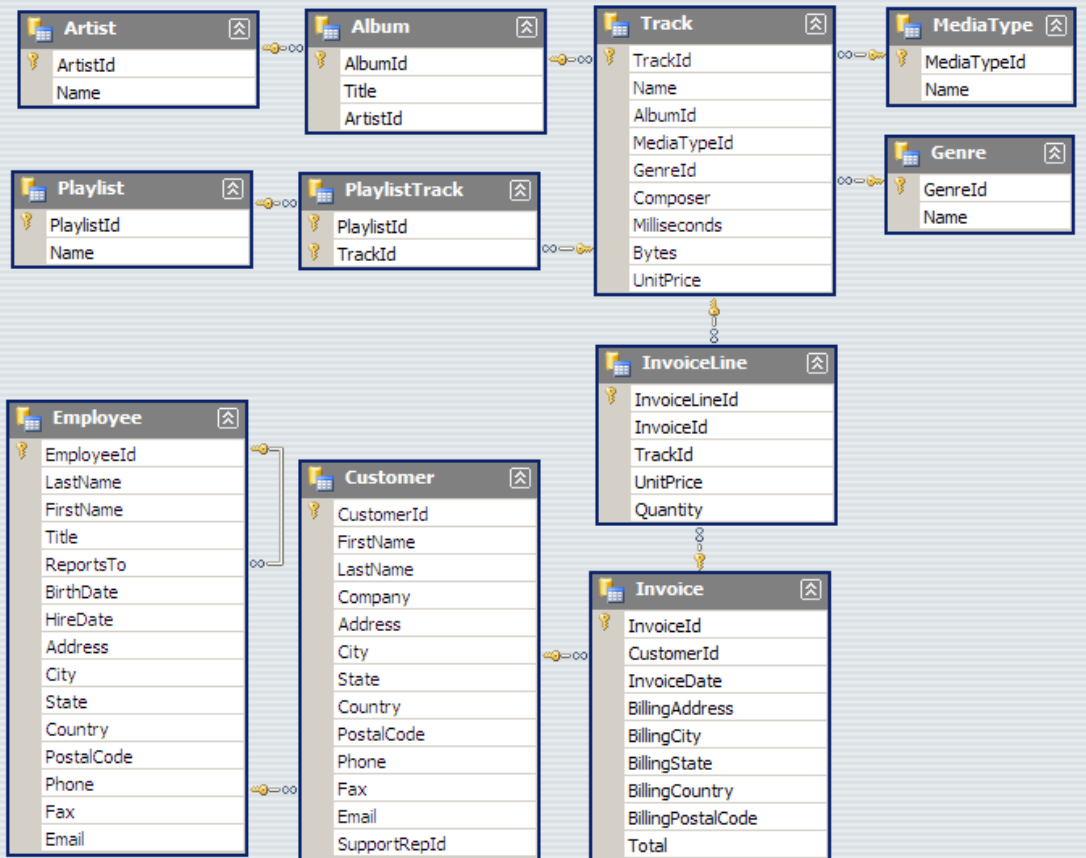


Test Databases

- Chinook Music Database
 - Used for examples
- Create chinook DB instance - .sql file on D2L

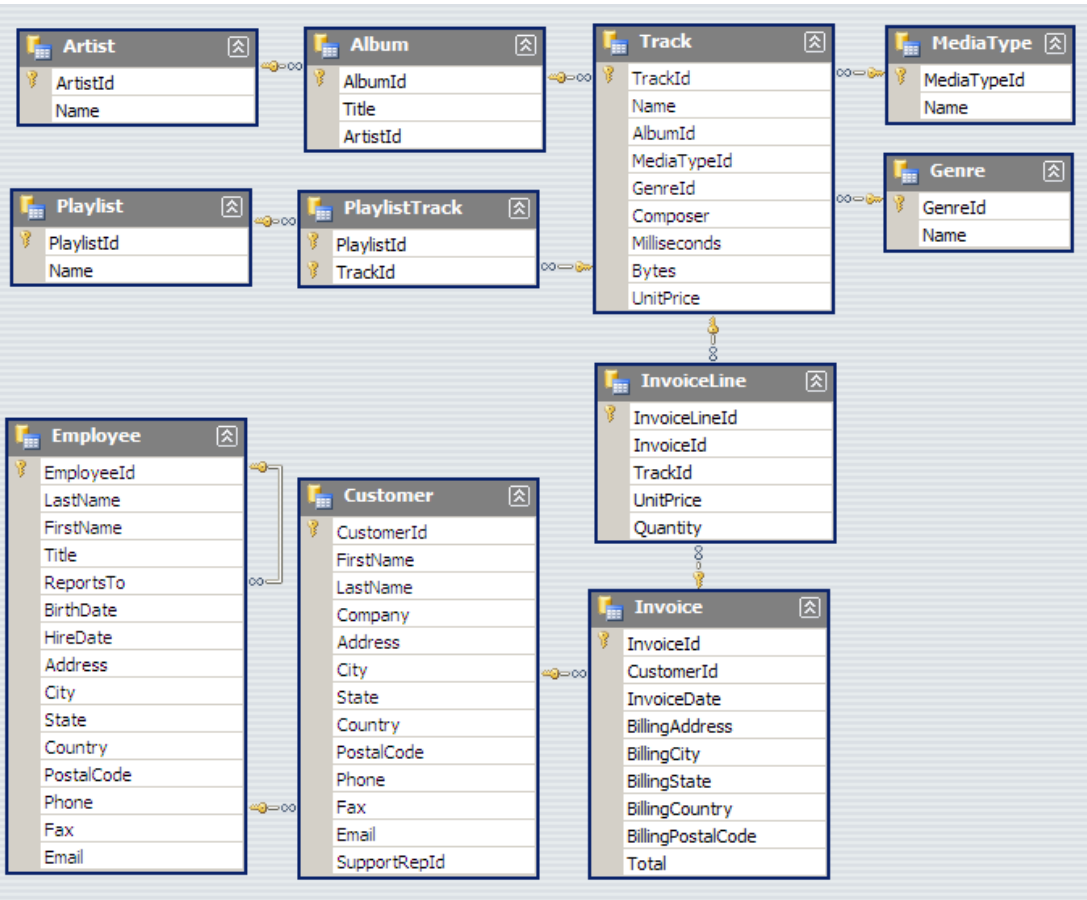
```
sqlite3 chinook.db < Chinook_Sqlite.sql
```

Practice



Provide a query showing customers who are not in the US

Practice

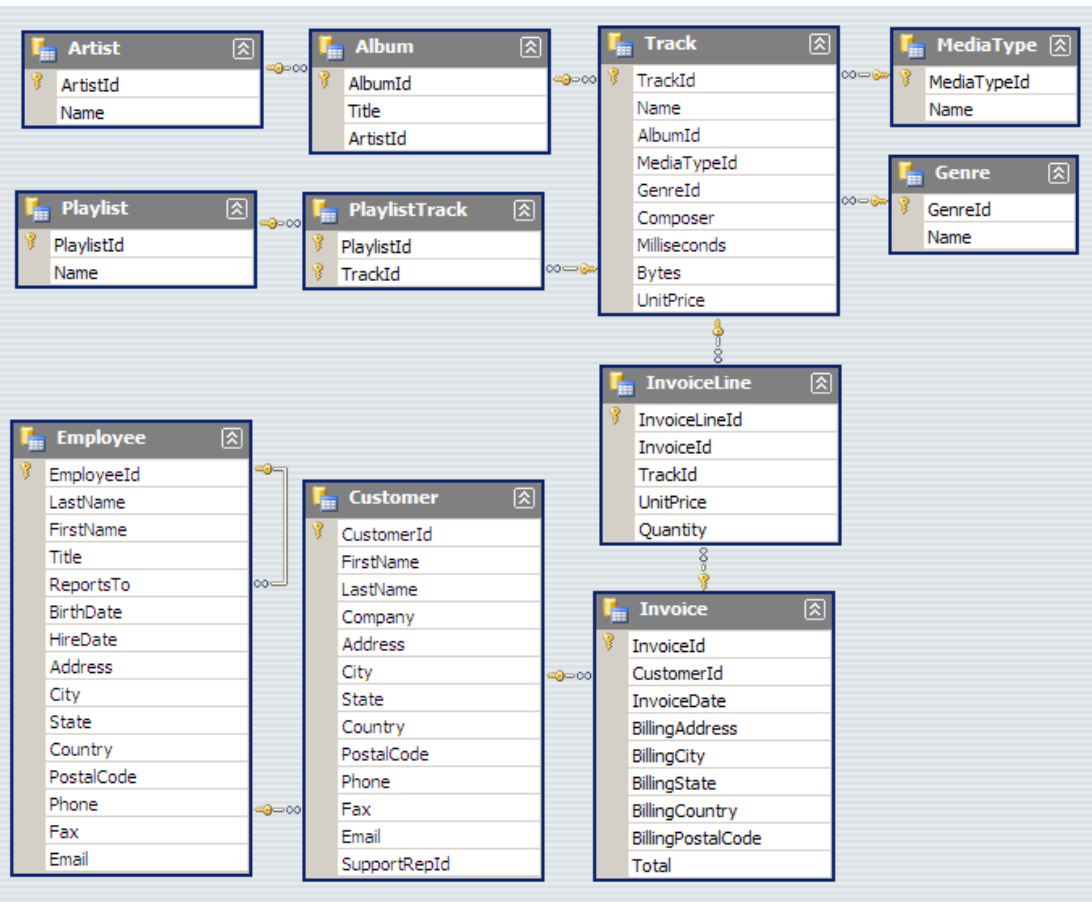


Answer:

```

SELECT FirstName, LastName, Country
FROM customer
WHERE NOT country = 'USA';
  
```

Practice



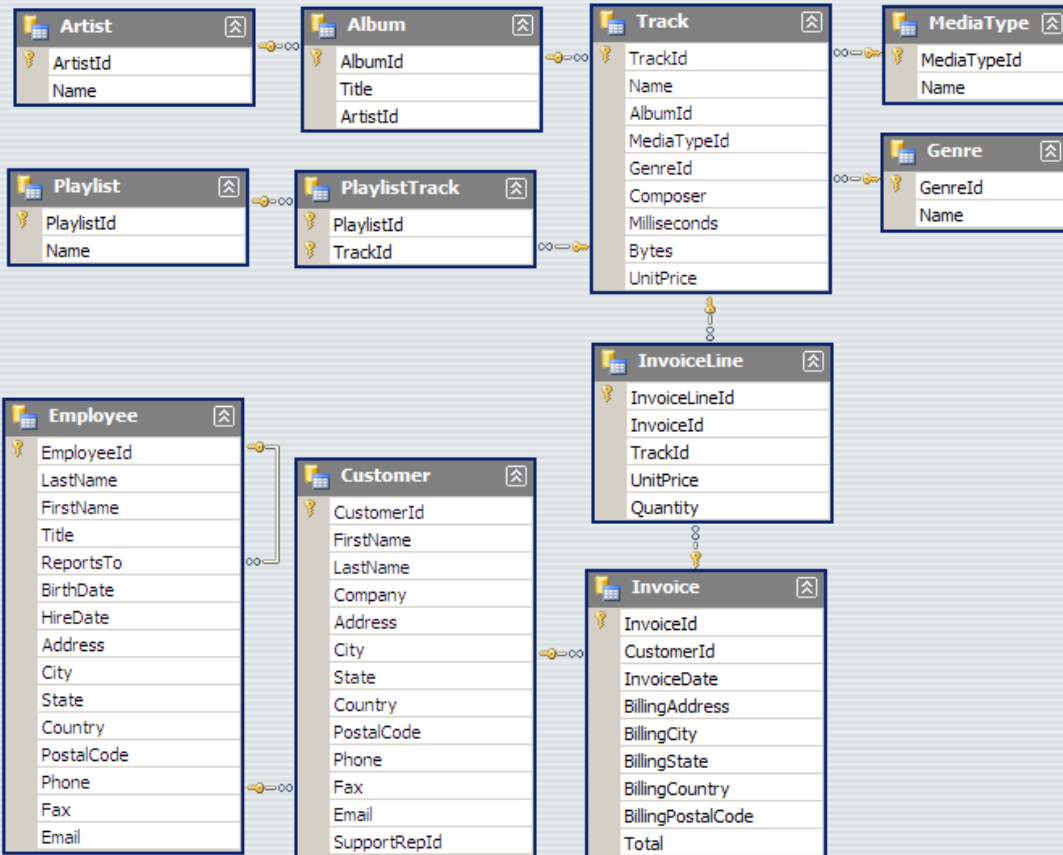
Provide a query showing invoices of customers who are from Brazil

Practice

Answer:

```
SELECT c.firstname, c.lastname,
i.invoiceid, i.invoicedate, i.billingcountry
FROM customer AS c INNER JOIN invoice
AS i ON c.customerid = i.customerid
WHERE c.country = 'Brazil';
```

Practice

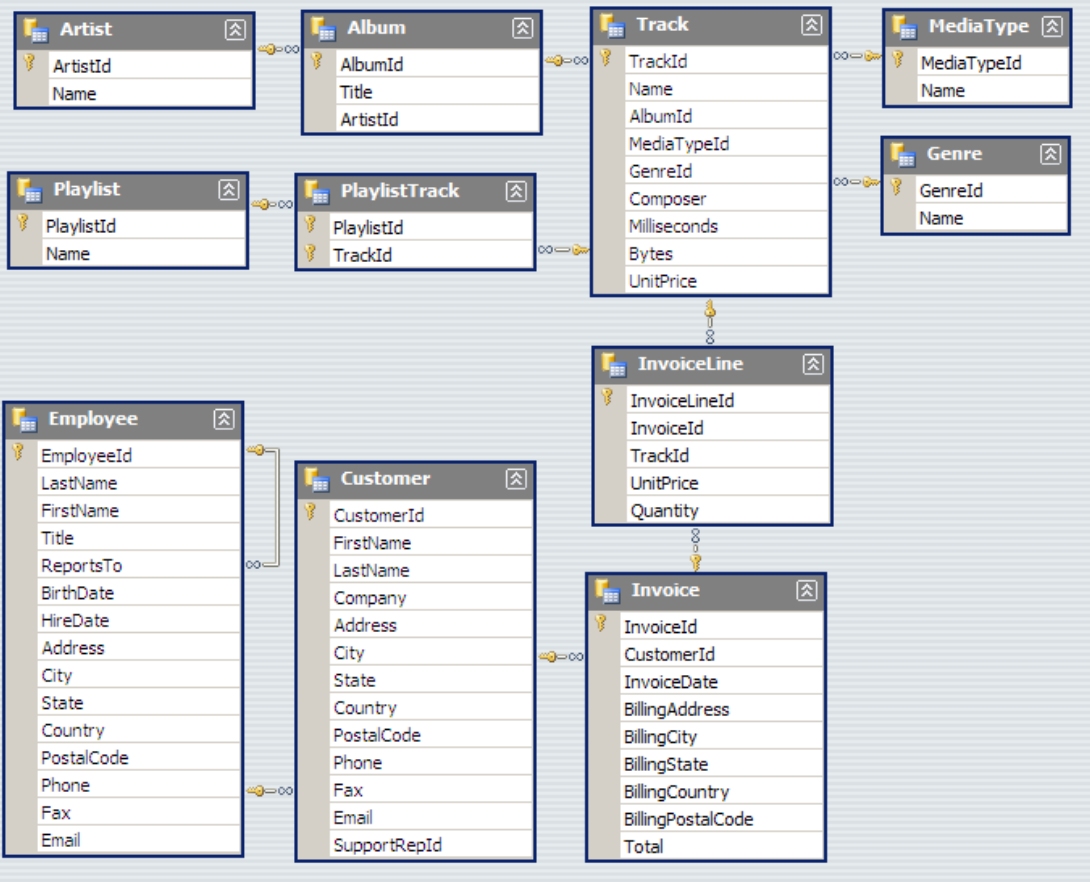


Answer:

```

SELECT c.firstname, c.lastname,
i.invoiceid, i.invoicedate, i.billingcountry
FROM customer AS c INNER JOIN invoice
AS i ON c.customerid = i.customerid
WHERE c.country = 'Brazil';
  
```


Practice



Answer:

```
SELECT c.firstname, c.lastname,
i.invoiceid, i.invoicedate, i.billingcountry
FROM customer AS c INNER JOIN invoice
AS i ON c.customerid = i.customerid
WHERE c.country = 'Brazil';
```

Practice

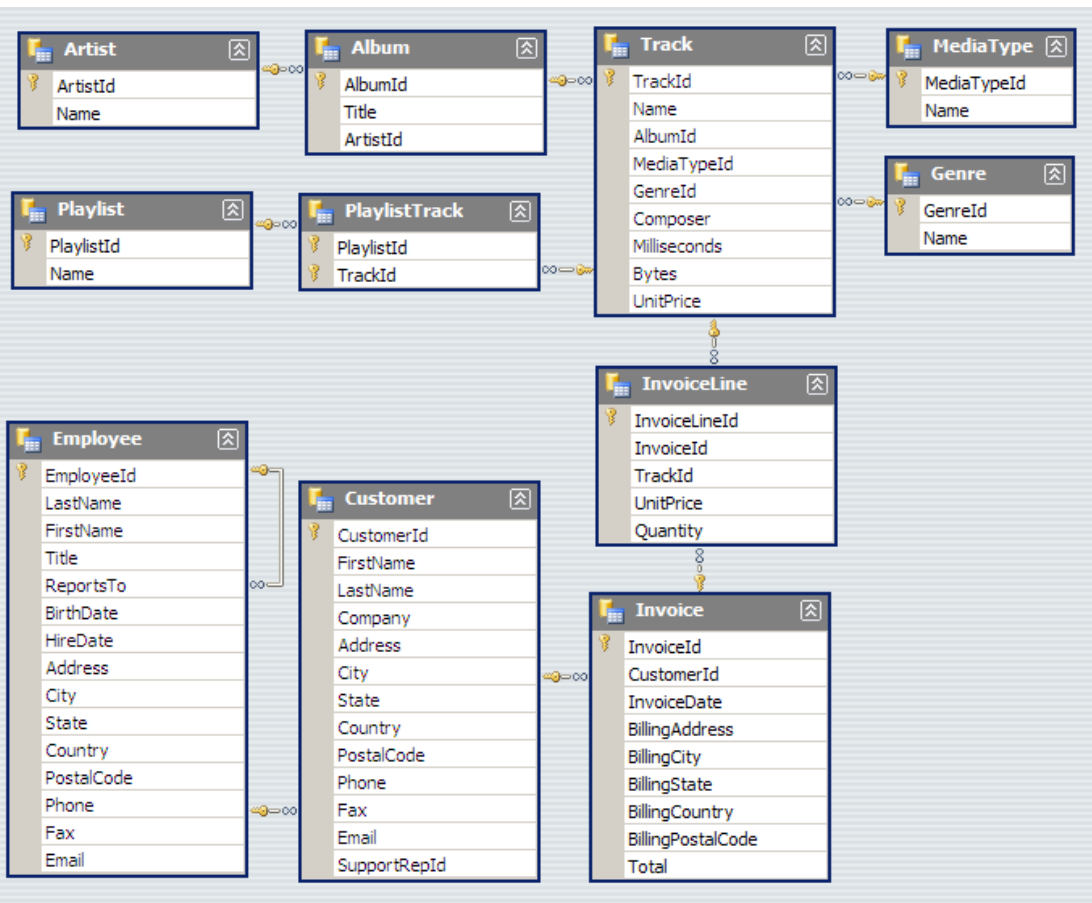


Answer:

```

SELECT c.firstname, c.lastname,
i.invoiceid, i.invoicedate, i.billingcountry
FROM customer AS c INNER JOIN invoice
AS i ON c.customerid = i.customerid
WHERE c.country = 'Brazil';
  
```


Practice

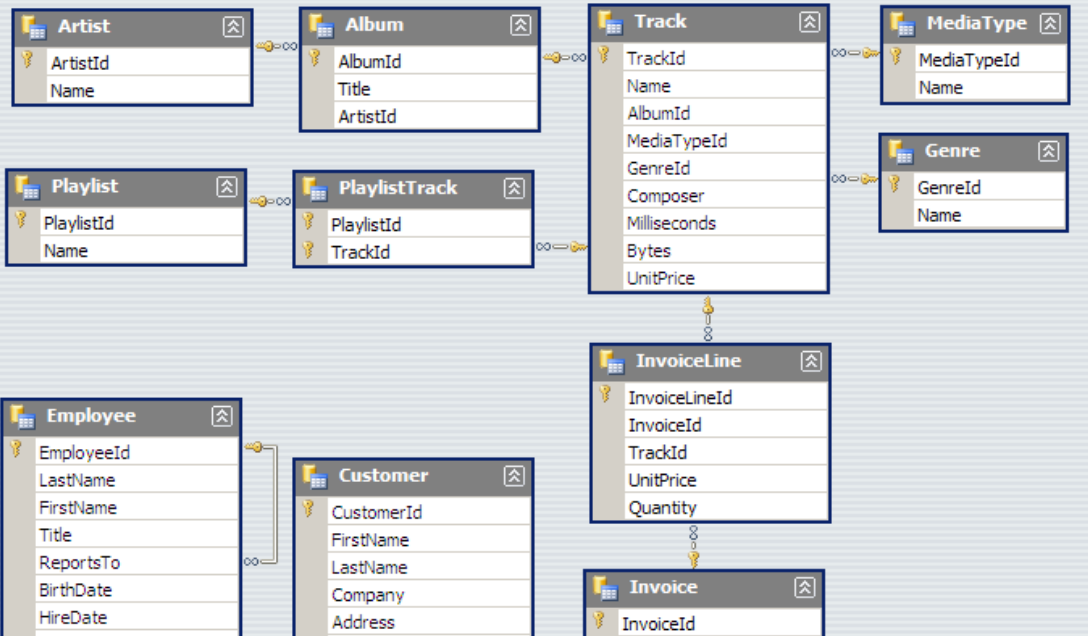


Provide a query showing each artist and their albums

Practice

Answer:

```
SELECT art.Name, alb.Title
FROM Artist AS art LEFT OUTER JOIN
Album AS alb
on art.ArtistId = alb.ArtistId;
```



Edson, DJ Marky & DJ Patife Featuring Fernanda Porto

Metallica
 Metallica
 Metallica
 Metallica
 Metallica
 Metallica
 Metallica
 Metallica
 Metallica
 Queen
 Queen
 Queen
 Kiss
 Kiss

Garage Inc. (Disc 1)
 Black Album
 Garage Inc. (Disc 2)
 Kill 'Em All
 Load
 Master Of Puppets
 ReLoad
 Ride The Lightning
 St. Anger
 ...And Justice For All
 Greatest Hits II
 Greatest Hits I
 News Of The World
 Greatest Kiss
 Unplugged [Live]

That's it for today

- Questions?