

# Spring 2024 CSE 380

1-23-2024

# Quiz Info

- We had our first quiz today
- Quiz 2 is moved to Feb 6<sup>th</sup>, the syllabus is updated
- I won't be returning quizzes, you can view a quiz in my office

# Review

# WHERE Clauses

- `SELECT ... FROM ... WHERE ...;`
- The where clause stipulates a test (predicate) that each row must pass to be returned in the select statement.

```
SELECT title, genre FROM Movies WHERE length  
< 120;
```

# WHERE Clauses

SELECT title, genre FROM Movies WHERE length  
< 120;

```
[sqlite> SELECT title, genre FROM Movies WHERE length < 120;
title          genre
-----
Wayne's World comedy
Back to the F  adventure
```

# Between

SELECT title, genre FROM Movies WHERE length  
**BETWEEN** 90 AND 130;

note: between is **inclusive of the bounds**

```
sqlite> SELECT title, genre FROM Movies WHERE length BETWEEN 90 AND 130;
title      genre
-----
Star Wars  sciFi
Wayne's Wo comedy
Back to th adventure
```

# IS and IS NOT

SELECT title, genre FROM Movies WHERE length  
IS NULL;

'IS' operator is only to be used to test for  
NULL values

```
sqlite> SELECT title, genre FROM Movies WHERE length is NULL;
title          genre
-----
Good Will Hunting drama
```

# IN

SELECT title, genre FROM Movies WHERE length  
IN(90,130);

- IN tests for inclusion in a static, parenthesized list

```
sqlite> SELECT title, genre FROM Movies WHERE length IN (90, 130);  
sqlite> █
```



# Comparing NULLS to Values

- Comparing any value (including NULL itself) with NULL yields UNKNOWN.
- A tuple is in a query answer if and only if the WHERE clause is TRUE (not FALSE or UNKNOWN).

# SQL Column Constraints

- Unique:
  - Used to specify a column whose values will always be unique for every row.
  - CREATE TABLE students (msu\_id TEXT UNIQUE,  
pid INTEGER UNIQUE,  
first\_name TEXT);

The DBMS will raise an error if the constraint is violated.

# NOT NULL

- Used to designate columns that may not be NULL

```
CREATE TABLE students (msu_id TEXT UNIQUE NOT  
NULL, section INTEGER NOT NULL);
```

- Raises error if a NULL value is inserted.

# Primary key

- Each table can have at most one PRIMARY KEY.
  - It is implicitly NOT NULL and UNIQUE
  - It is used to specify a specific column in the table

```
CREATE TABLE students (msu_id TEXT PRIMARY KEY,  
                        grade REAL);
```

# Joins

- Often, queries need to combine (join) rows from multiple tables.
- Joins specify which rows get merged with rows from a second table.
- As there are many possible methods to match rows together, there are many types of joins.
- You will need to know (memorize) the different joins.
- Annoying, but very useful knowledge.

# Joins

professors

id	name
esfahanian	Abdol
mariani	James
dyksen	Wayne
NULL	demo-prof

```
CREATE TABLE professors (id TEXT, name TEXT);
INSERT INTO professors VALUES('esfahanian','Abdol');
INSERT INTO professors VALUES('mariani','James');
INSERT INTO professors VALUES('dyksen','Wayne');
INSERT INTO professors VALUES(NULL, 'demo-student');
```

associates

responder_id	associate_id
mariani	dyksen
mariani	esfahanian
dyksen	mariani
esfahanian	dyksen

```
CREATE TABLE associates (responder_id TEXT, associate_id TEXT);
INSERT INTO associates VALUES('mariani','dyksen');
INSERT INTO associates VALUES('mariani','esfahanian');
INSERT INTO associates VALUES('dyksen','mariani');
INSERT INTO associates VALUES('esfahanian','dyksen');
```

# Inner Join


- INNER JOIN is like the CROSS JOIN, except only rows that match a predicate are allowed. Frequently that predicate includes primary keys to match associated data

```
SELECT * FROM professors INNER JOIN  
associates  
    ON id = responder_id;
```

- You can equivalently do an implicit inner join and use WHERE to filter rows:

```
SELECT * FROM students, associates  
    WHERE id = responder_id;
```

# Inner Join



Returned rows are those that match both tables (shaded).



# Inner Join

- `SELECT * FROM professors INNER JOIN associates ON id = responder_id;`


```
sqlite> SELECT * FROM professors INNER JOIN associates ON id = responder_id;
```

id	name	responder_id	associate_id
esfahanian	Abdol	esfahanian	dyksen
mariani	James	mariani	dyksen
mariani	James	mariani	esfahanian
dyksen	Wayne	dyksen	mariani

# Left Outer Join

- LEFT OUTER JOIN includes all the rows to the left of the JOIN keyword, where there is a match, it add the matching column data to rows. If there isn't a match NULL values are used instead.
- `SELECT * FROM professors LEFT OUTER JOIN associates  
ON id = responder_id;`
- Note: order matters!!!
- `SELECT * FROM associates LEFT OUTER JOIN professors  
ON id = responder_id;`

# Left Outer Join



# Left Outer Join

SELECT \* FROM professors LEFT OUTER JOIN associates  
ON id = responder\_id;

```
[sqlite> SELECT * FROM professors LEFT OUTER JOIN associates ON id=responder_id;
```

id	name	responder_id	associate_id
esfahanian	Abdol	esfahanian	dyksen
mariani	James	mariani	dyksen
mariani	James	mariani	esfahanian
dyksen	Wayne	dyksen	mariani
	demo-stude		

# New Material

# SQL Union

- UNION combines the result sets of two queries.
- Column data types in the two queries must match.
- UNION combines by column position rather than column name.
- If we wanted all the names of Artists and Albums:

```
SELECT Artist.Name  
FROM Artist  
UNION  
SELECT Album.Title  
FROM Album;
```

# SQL Union

```
sqlite> SELECT Artist.Name FROM Artist WHERE Artist.Name IN (SELECT Album.Title FROM Album) ORDER BY Artist.Name DESC LIMIT 10;  
Name  
-----  
Van Halen  
The Doors  
Temple of  
Raul Seixa  
Pearl Jam  
Olodum  
Iron Maide  
Body Count  
Black Sabb  
Audioslave
```

```
SELECT Artist.Name  
FROM Artist  
UNION  
SELECT Album.Title  
FROM Album;
```

# SQL Union All

- UNION performs a DISTINCT on the result set, eliminating any duplicate rows.
- UNION ALL does not remove duplicates, and it therefore faster to perform than UNION.
- If I wanted the Artists and Album names, but duplicates were okay:

```
SELECT Artist.Name  
FROM Artist  
UNION ALL  
SELECT Album.Title  
FROM Album;
```



# When is a Left Outer Join the Same as an Inner Join?

- Never
- When the left rows always have a match
- When the left rows don't have NULLs
- Always

# SQL Subqueries

- A subquery is a SQL query within a query.
- Subqueries are nested queries that provide data to the enclosing query.
- Subqueries can return individual values or a list of records
- Subqueries must be enclosed with parenthesis
- If I wanted all the Artists who released an Album with just their own name:

```
SELECT Artist.Name  
FROM Artist  
WHERE Artist.Name IN (SELECT Album.Title FROM  
Album);
```

# SQL Subqueries

- Subqueries often used to perform tests for set membership, make set comparisons, and determine set cardinality

**SELECT name FROM courses WHERE  
semester = 'Fall' AND semester = 'Spring';**

**SELECT name FROM courses WHERE semester  
= 'Fall' AND name in (SELECT name FROM  
courses WHERE semester = 'Spring');**

name	semester
CSE 480	Fall
CSE 480	Spring
CSE 498	Fall
CSE 498	Spring
CSE 422	Fall
CSE 476	Fall

```
[sqlite> SELECT name FROM courses WHERE semester = 'Fall' AND semester = 'Spring';  
sqlite> 
```

```
[sqlite> SELECT name FROM courses WHERE semester = 'Fall' AND name IN (SELECT name FR  
OM courses WHERE semester = 'Spring');  
name  
-----  
CSE 480  
CSE 498  
sqlite> 
```

# Rest of SQL SELECT

- MIN MAX
  - Select min return the minimum value for a column
  - Select max return the maximum value for a column

SELECT first\_name, MIN(siblings) FROM students;

```
[sqlite> SELECT first_name, MIN(siblings) FROM students;
first_name  MIN(siblings)
-----
Ishita      0
```

SELECT first\_name, MAX(siblings) FROM students;

```
[sqlite> SELECT first_name, MAX(siblings) FROM students;
first_name  MAX(siblings)
-----
Odon        8
```

# COUNT, SUM, TOTAL, AVG

- COUNT returns a count of non-null values
- SUM returns the sum of the data (NULL if all rows are NULL)
- TOTAL returns total sum (NULL is zero)
- AVG returns average of data values

```
[sqlite> SELECT COUNT(siblings) FROM students;  
COUNT(siblings)  
-----  
110
```

```
[sqlite> SELECT SUM(siblings) FROM students;  
SUM(siblings)  
-----  
180
```

```
[sqlite> SELECT TOTAL(siblings) FROM students;  
TOTAL(siblings)  
-----  
180.0
```

```
[sqlite> SELECT AVG(siblings) FROM students;  
AVG(siblings)  
-----  
1.63636363636364
```

# DISTINCT

- SELECT DISTINCT returns only distinct (different) values.
- SELECT DISTINCT eliminates duplicate records from the results.
- DISTINCT can be used with aggregates: COUNT, AVG, MAX, etc.
- DISTINCT operates on a single column. DISTINCT for multiple columns is not supported.

SELECT DISTINCT siblings FROM students;

```
[sqlite> SELECT DISTINCT siblings FROM students;
siblings
-----
1
0
2
3
5
8
4
6
```

SELECT count(DISTINCT siblings) FROM students;

```
[sqlite> SELECT COUNT(DISTINCT siblings) FROM students;
COUNT(DISTINCT siblings)
-----
8
```

# GROUP BY

- The GROUP BY clause groups records into summary rows.
- GROUP BY returns one records for each group.
- GROUP BY typically also involves aggregates: COUNT, MAX, SUM, AVG, etc.
- GROUP BY can group by one or more columns.

SELECT count(), siblings FROM students GROUP BY siblings;

```
[sqlite> SELECT count(), siblings FROM students GROUP BY siblings;
count()      siblings
-----
1
13           0
43           1
38           2
11           3
1            4
2            5
1            6
1            8
```

# LIMIT

- Last clause of SELECT statement.
- Limits the number of rows to the value (or less).
- Often useful with ORDER BY to get extremes.

SELECT first\_name, siblings FROM students  
ORDER BY siblings DESC LIMIT 10;

```
sqlite> SELECT first_name, siblings FROM students ORDER BY siblings DESC LIMIT 10;
first_name  siblings
-----
Odon        8
Jimmy       6
Yunpeng     5
Frederick   5
Andrew      4
Emily       3
Ajuisiwon   3
Finn        3
Brandon     3
Gregory     3
```



# All the parts of a SELECT query

- SELECT column-names
- FROM table-name
- WHERE condition
- GROUP BY column-names
- HAVING condition
- ORDER BY column-names
- LIMIT max-rows
- ;

# CRUD

- Four fundamental operations that apply to any database are:
  - Read the data -- *SELECT*
  - Insert new data -- *INSERT*
  - Update existing data -- *UPDATE*
  - Remove data -- *DELETE*
- Collectively these are referred to as **CRUD** (Create, Read, Update, Delete).

# INSERT INTO

- The INSERT INTO statement is used to add new data to a database.
- The INSERT INTO statement adds a new record to a table.
- INSERT INTO can contain values for some or all of its columns.
- INSERT INTO can be combined with a SELECT to insert records.

```
INSERT INTO spelling_team, (first_name, spelling)  
VALUES ('James', 10);
```

# INSERT Multiple Values

- You can insert multiple rows with a single INSERT statement:

```
INSERT INTO spelling_team (first_name, spellings)
VALUES ('James', 10),
      (Abigail', 9);
```

# INSERT INTO SELECT

- You can also use a SELECT clause to generate the needed rows, but you need to return the correct column types and order.

```
INSERT INTO spelling_team (first_name, spelling)
SELECT students.first_name, students.spelling
FROM students ORDER BY spelling DESC LIMIT 10;
```

```
[sqlite> SELECT * FROM spelling_team;
first_name  spelling
-----
Oscar       12
Yujin       10
Shafkat     10
Rohit       10
Matthew     10
Marla       10
Ishita      10
Elio        10
David       10
Brandon     10
```

# UPDATE

- The UPDATE statement updates data values in a database.
- UPDATE can update one or more records in a table.
- Use the WHERE clause to UPDATE only specific records.

```
UPDATE students SET spelling = 0 WHERE spelling >10;
```

```
[sqlite> SELECT * FROM spelling_team;
first_name  spelling
-----
Yujin       10
Shafkat     10
Rohit       10
Matthew     10
Marla       10
Ishita      10
Elio        10
David       10
Brandon     10
Antonio     10
```

# UPDATE

- The UPDATE statement updates data values in a database.
- UPDATE can update one or more records in a table.
- Use the WHERE clause to UPDATE only specific records.

```
UPDATE students SET spelling = 0 WHERE spelling >10;
```

```
UPDATE students SET fav_word = fav_word || '!';
```

|| is concatenate in SQL

# DELETE

- DELETE permanently removes records from a table.
- DELETE can delete one or more records in a table.
- Use the WHERE clause to DELETE only specific records.

```
DELETE FROM spelling_team; -- removes all rows
```

```
DELETE FROM spelling_team WHERE got_answer_right = 'False';  
--removes some rows
```



# AUTOINCREMENT


- SQLite AUTOINCREMENT is a keyword used for auto incrementing a value of a field in the table. We can auto increment a field value by using AUTOINCREMENT keyword when creating a table with specific column name to auto incrementing it.
- The keyword AUTOINCREMENT can be used with a INTEGER PRIMARY KEY field only.
- CREATE TABLE table\_name ( column1 INTEGER PRIMARY KEY AUTOINCREMENT, column2 datatype, column3 datatype, ..... columnN datatype);

# Should you use it?

- From <https://www.sqlite.org/autoinc.html>:  
"The AUTOINCREMENT keyword imposes extra CPU, memory, disk space, and disk I/O overhead and should be avoided if not strictly needed. It is usually not needed."
- If a column is INTEGER PRIMARY KEY, it already will autofill a unique value if a value isn't provided.
  - The only difference is AUTOINCREMENT guarantees monotonically increasing values, instead of just unique.

## Practice

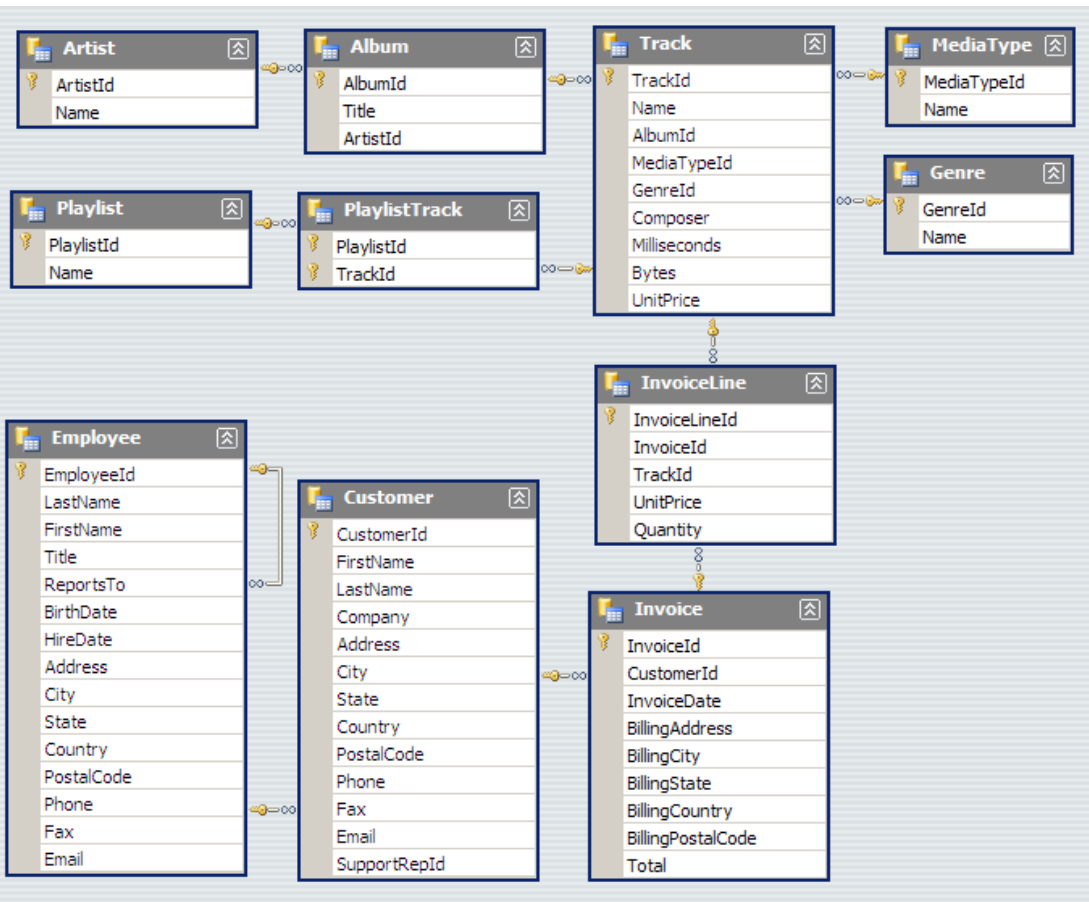
# Chinook Dataset



# Test Databases


- Chinook Music Database
  - Used for examples
- Create chinook DB instance - .sql file on D2L

```
sqlite3 chinook.db < Chinook_Sqlite.sql
```



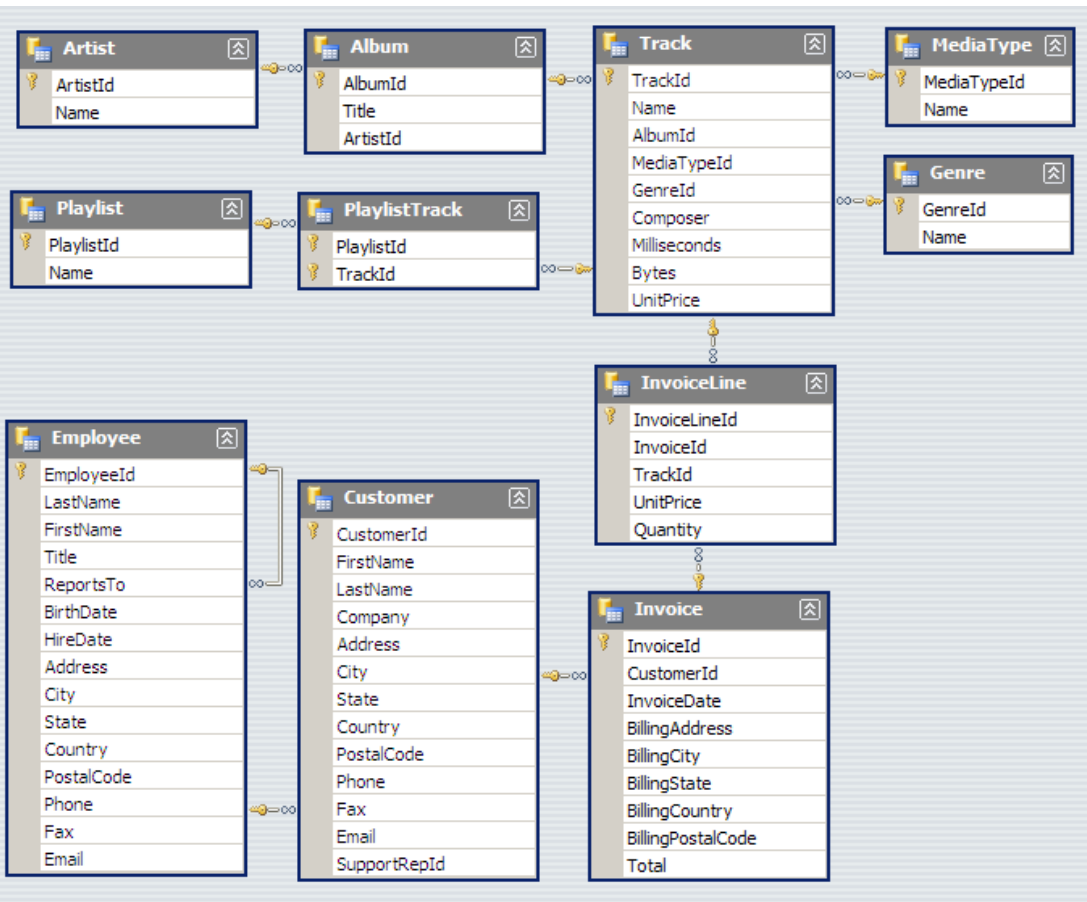
## Practice

Provide a query that shows the total sales per country in order from most to least sales



## Practice

Provide a query that shows the number of customers assigned to each sales agent



## Practice

Provide a query that shows the total number of tracks in each playlist, and includes the playlist name in the result

# That's it for today

- Questions?