

# Spring 2024 CSE 380

2-1-2024

# Quiz Info

- Quiz 2 is next Tuesday, February 6th

# Review

# Foreign Keys

- We discussed primary keys, which are columns that uniquely identify each row.
- However, often our tables will have columns that are meant to match up with columns in a different table.
- We want to add a constraint on those columns that there must be an associated row in a foreign (other) table.
  - NULLs are okay

# More Table Constraints


tracks

| id | name                       | artist_id |
|----|----------------------------|-----------|
| 1  | Let it be                  | 1         |
| 2  | Penny Lane                 | 1         |
| 3  | Yellow Submarine           | 1         |
| 4  | Hedwig's Theme             | 2         |
| 5  | Jingle Bell Rock           | NULL      |
| 6  | The Devil Is A Patient Man | 3         |

tracks.artist\_id is a  
foreign key  
referencing artists.id

artists

| id | name          |
|----|---------------|
| 1  | The Beatles   |
| 2  | John Williams |
| 3  | CRUD          |



# Example

```
PRAGMA foreign_keys = ON;  
INSERT INTO artists (id, name) VALUES (1, 'Beatles');  
INSERT INTO tracks (id, name, artist_id) VALUES  
    (1, 'Jingle Bell Rock', NULL), -- OKAY  
    (2, 'Let it be', 1), -- Okay  
    (3, 'Jurassic Park', 2); -- ERROR: no matching key
```

# Parameterized Queries

```
conn = sqlite3.connect(":memory:")  
conn.execute("CREATE TABLE students (name TEXT, age INTEGER);")  
conn.execute("INSERT INTO students VALUES ('James', 30);")
```

What if we wanted to add a python integer?

```
Steve_age = 23
```

```
conn.execute( "INSERT INTO students VALUES ('Steve', " + str(steve_age) + ");" )
```

# Parameterized Queries

- Used to pass python objects into queries without needing to manually convert to strings.

```
Steve_age = 23
```

```
conn.execute("INSERT INTO students VALUES ('Steve', ?);", (steve_age,))
```

```
row = ('Tim', 45)
```

```
conn.execute("INSERT INTO students VALUES (?, ?);", row)
```



# Protection by Parameterized Query

- Parameterized queries will automatically escape the input and ensure that the value passed in is stored as that value.

```
name = "Robert'); DROP TABLE students; --"
```

```
conn.execute("INSERT INTO students VALUES (?)", (name,))
```

- The string name will be stored, in its entirety, and the single quote will be escaped to stop it from harming the database.

# Python and SQLite

# Example Code

- This will be posted online

```
import sqlite3
conn = sqlite3.connect("test.db")
curr = conn.cursor()

curr.execute("DROP TABLE IF EXISTS students")
curr.execute("CREATE TABLE students (col1 INTEGER, col2 TEXT, col3 REAL);")

curr.execute("INSERT INTO students VALUES (3, 'hi', 4.5);")

multiple_records = [(7842, 'string with spaces', 3.0), (7, 'look a null', None)]
curr.executemany("INSERT INTO students VALUES (?, ?, ?);", multiple_records)

curr.execute("SELECT col1, col2, col3 FROM students ORDER BY col1;")
result_list = curr.fetchall() #fetchone(), fetchmany(3)

expected = [(3, 'hi', 4.5), (7, 'look a null', None), (7842, 'string with spaces', 3.0)]

print("expected:", expected)
print("actual: ", result_list)
assert expected == result_list
```

# New Material

# Functions

- Functions take 0 or more arguments and return one result. Here are some more:
  - `abs(x)` - returns the absolute value of its argument
  - `lower(x)` - returns a copy of a string in lower case
  - `upper(x)` - guess what this does
  - `random(x)` - returns an random, signed 64-bit integer
  - `typeof(x)` - returns the type ("null", "integer", "real", "text", "blob") of x
  - `round(x, y)` - returns a floating-point value x rounds to y digits after the decimal point
  - `trim(x, y)` - returns a copy of string x with y characters removed from each end

# User Defined Functions

- What if we wanted to add a custom function to SQLite?
  - title(x) - returns the string x in "Title Case" (the first letter of each word is capitalized).
- We first write a function (in Python) that does what we want:

```
def make_title(x):  
    return x.title()
```

- Then we need to tell our connection to the database that this function exists:

```
conn.create_function("title", 1, make_title)
```

- Now we can use it in SQL queries:

```
curr.execute("SELECT title(name) FROM students;")
```

# create\_function

- The create\_function method on the connection object takes three parameters:
  - name - a string giving the name the function will be called by in SQL
  - num\_params - the number of parameters the function requires
  - You can overload functions with different number of arguments
  - func - the function (in Python) to be called
- The function can return any of the SQLite supported types (bytes, str, int, float, and None)

# Aggregate Functions

- Aggregators (aggregate functions) take 0 or more values and return some form of summary or those values. When seen many before:
  - `count(x)` - returns the number of items in x
  - `max(x)` - returns the largest value in x
  - ...
- You can also create a custom aggregator with the Python `sqlite3` module
  - Lets say we wanted an aggregate function called `char_length` that tallied the number of characters in the strings passed to it

- Example Usage:

```
CREATE TABLE x (col TEXT);
```

```
INSERT INTO x VALUES ('hi'), ('bye');
```

```
SELECT char_length(col) FROM x; -- returns 5
```



# Custom Aggregator

```
class CharCounter:
    def __init__(self):
        self.char_count = 0
    def step(self, value):
        self.char_count += len(value)
    def finalize(self):
        return self.char_count
conn.create_aggregate("char_length", 1, CharCounter)
```

# create\_aggregate

- Just like `create_function`, the `create_aggregate` method of the SQLite3 Connection class has three arguments:
  - `name` - the SQL name to call the aggregate function by
  - `num_params` - the number of parameters for the aggregate function
  - `aggregate_class` - the Python class to pass data to
- The aggregate class must have three methods:
  - `__init__(self)` - the init method that initializes the class (usually an attribute to 0)
  - `step(self, value)` - a method that is called for every values to be aggregated
  - `finalize(self)` - a method that returns the value for the aggregate function/class

# Storing Custom Types

- Often you want to store different types of data in a database. The database can only hold a few types of data (INTEGER, REAL, TEXT, BLOB, NULL), but you can use adapters and converters to change objects into these few types and back.
- Adapter - A function to convert a Python type into one of SQL's supported types (string or bytestring)
- Converter - A function to convert a bytestring (BLOB) to a Python type.
- We've can use these whenever we store Python datetime objects in SQLite.

# Custom Python Type

- Lets say we wanted to represent colors in our database. We have a Color class in Python:

```
class Color:
    def __init__(self, r, g, b):
        self.r, self.g, self.b = r, g, b
    def __repr__(self): return "Color({}, {}, {})".format(
        self.r, self.g, self.b)
```

- Then we need two functions (an adapter and a converter) to convert to a form that can be put into a Sqlite database and returned to being an Python object.

# Adapter and Converter

```
def adapt_color(color):  
    return "{};{};{}".format(color.r, color.g, color.b)  
  
def convert_color(bytestring):  
    as_str = bytestring.decode('ascii')  
    r, g, b = [float(x) for x in as_str.split(';')]  
    return Color(r, g, b)  
  
sqlite3.register_adapter(Color, adapt_color)  
sqlite3.register_converter("COLOR", convert_color)
```

# Python Types

- Adapters take a custom python object (like a "Color" object), and return a string containing the data from that object.
- Converters take a bytes object (python binary string) and should return a new python object instance.
- You can convert a bytes object to the more familiar python strings with:
  - `as_str = as_bytes.decode('ascii')`
  - Now `as_str` holds a python string that you can split and do other things with.

# Using with PARSE\_DECLTYPES

```
c = Color(24, 69, 59)
conn = sqlite3.connect(":memory:", detect_types=sqlite3.PARSE_DECLTYPES)
curr = conn.cursor()
curr.execute("CREATE TABLE test(col COLOR);")
curr.execute("INSERT INTO test(col) VALUES(?);", (c,))
curr.execute("SELECT col FROM test;")
row = curr.fetchall()
row = next(res)
print("with declared types:", row[0])
```

"detect\_types=sqlite3.PARSE\_DECLTYPES" instructs SQLite to try to convert values inserted into a column with a custom type to the associated Python type.

# Views

- Views are named SELECT Statement

| students | id | name    | log | student_id | time_inserted |
|----------|----|---------|-----|------------|---------------|
|          | 1  | James   |     | 1          | 12:40:00      |
|          | 2  | Abigail |     | 2          | 14:37:34      |

```
CREATE VIEW name_and_time AS
  SELECT students.name, log.time_inserted FROM students
    INNER JOIN log ON students.id = log.student_id;
```

| name_and_time | name    | time_inserted |
|---------------|---------|---------------|
|               | James   | 12:40:00      |
|               | Abigail | 14:37:34      |



# Views Cont.

- A view is a virtual table composed of the result-set of a select statement.
- A view has rows and columns (just like a table), but instead of holding data itself, it just points to data held in other tables.
- You can treat a view like a table and do select statements on it,
  - But views cannot be modified (no insert, update or delete)
- You can use a view like a temporary table. You can build the view from multiple other tables with joins and/or filter certain rows with WHERE and so on.
- If the view's SELECT statement changes (modifications where made to the SELECT's underlying tables), the view automatically changes too.

# View Example

```
CREATE VIEW artist_and_albums AS
    SELECT Artist.Name, Album.Title, Album.AlbumId
    FROM Artist INNER JOIN Album ON Artist.ArtistId = Album.ArtistId;
CREATE VIEW good_artists_and_albums AS
    SELECT * FROM artist_and_album WHERE NOT Name LIKE '%Taylor Swift%';
CREATE VIEW good_tracks AS
    SELECT Tracks.* FROM Tracks INNER JOIN good_artists_and_albums
    ON Tracks.AlbumId = good_artists_and_albums.AlbumId;
```

# When Should you use Views?

1. When you want a read-only table made from data in other tables.
  2. To hide data complexity (for instance, complicated joins).
  3. Customizing the data, using functions and group by, to present the data in a new form.
  4. Security and privacy
- However, views are as slow as the query that creates them. And, views of views are even slower as they are basically nested SELECT subqueries.

# View Lifespan

- A view lives only as long as the tables in its select statement.
  - If the tables in the select statement are dropped, the view is also dropped.
- You can use the TEMPORARY keyword to indicate that the view should automatically be dropped when the current database connection is closed (meaning the view isn't persistent across connections).

CREATE TEMPORARY VIEW something AS SELECT \* FROM students;

You can also remove a view manually with DROP VIEW:

DROP VIEW something;

# Indices

**Note: Indices  
not on Quiz 2**

- What is an index?
  - Like the index in the back of a book (remember those?), an index helps you find information quickly.
  - Without an index, you would have to go through the entire book to find a particular topic.
    - But with an index, you can look up the topic in a sorted area, and then jump to the topic you care about.
  - Indices have no external effect on a database; they do not affect any of the other queries.
    - But they do have a huge effect on performance.
  - By default, every table is indexed according to its integer primary key. If there is no primary key, SQLite makes a hidden one called "rowid".

# Create Index

CREATE INDEX index\_name ON students (name);

- You can add an index to a database with the above command

CREATE INDEX index\_name students (id, name);

- You can have multiple columns to sort your index on.

DROP INDEX index\_name;

DROP INDEX IF EXISTS index\_name;

- As always IF [NOT] EXISTS makes the CREATE/DROP a no-op if an table/view/index with that name already exists

# When Should you Create an Index?

- When you need to often look up rows, but not according to their primary key.
- An index on some columns in a table allows it to quickly get matching rows when doing a lookup on those columns.
- But, indices aren't always good, modifications to a table (UPDATE, INSERT, DELETE) all need to adjust the indices on the table, leading to slower performance for those operations.
- Also, indices take up memory, so large/multiple indices can take up valuable space in your database.

# Query Planner

- The following content is taken from <https://www.sqlite.org/queryplanner.html>.
- SQL is a declarative language, not a procedural language, meaning you write what you want to do, not how to do it.
- It is the job of the SQL implementation to interpret your declaration into actual steps to yield the answer.
- Most of the time the implementation can do this efficiently, but sometimes it needs you to create indices to make the operations you need done faster.



## fruitsforsale

```
CREATE TABLE
fruitsforsale (
    fruit TEXT,
    state TEXT,
    price REAL
);
INSERT INTO ...
```

| rowid | fruit      | state | price |
|-------|------------|-------|-------|
| 1     | Orange     | FL    | 0.85  |
| 2     | Apple      | NC    | 0.45  |
| 4     | Peach      | SC    | 0.60  |
| 5     | Grape      | CA    | 0.80  |
| 18    | Lemon      | FL    | 1.25  |
| 19    | Strawberry | NC    | 2.45  |
| 23    | Orange     | CA    | 1.05  |

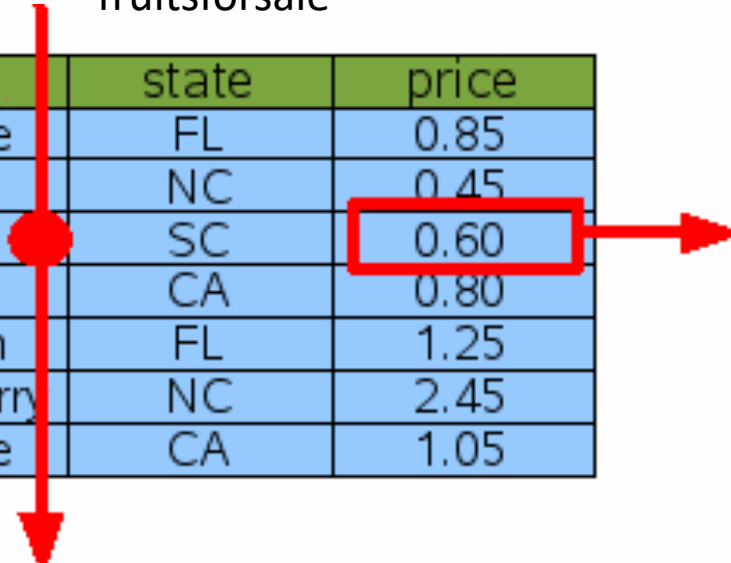
We didn't specify a INTEGER PRIMARY KEY, so a new column "rowid" was made for us as the INTEGER PRIMARY KEY.

The table is always ordered according to the key.

```
SELECT price FROM  
fruitsforsale WHERE  
fruit = 'Peach';
```

fruitsforsale

| rowid | fruit      | state | price |
|-------|------------|-------|-------|
| 1     | Orange     | FL    | 0.85  |
| 2     | Apple      | NC    | 0.45  |
| 4     | Peach      | SC    | 0.60  |
| 5     | Grape      | CA    | 0.80  |
| 18    | Lemon      | FL    | 1.25  |
| 19    | Strawberry | NC    | 2.45  |
| 23    | Orange     | CA    | 1.05  |



Without an index on fruit, we have to do a full table scan, meaning we must examine every row and check if fruit is equal to 'Peach'.

This is very slow if you have thousands or millions of rows.

## What is the Big O Notation for a Simple SELECT Statement? (n is number of rows)

- $n$
- $\log(n)$
- $\log(n) + \log(n)$
- Uhhhh...

```
CREATE INDEX  
fruit_index ON  
fruitsforsale(fruit);
```

Now we have a index ordered  
by fruit so we can quickly find  
rows with a certain fruit value.

fruit\_index

| fruit      | rowid |
|------------|-------|
| Apple      | 2     |
| Grape      | 5     |
| Lemon      | 18    |
| Orange     | 1     |
| Orange     | 23    |
| Peach      | 4     |
| Strawberry | 19    |

```
SELECT price FROM fruitsforsale WHERE fruit = 'Peach';
```

fruit\_index

| fruit      | rowid |
|------------|-------|
| Apple      | 2     |
| Grape      | 5     |
| Lemon      | 18    |
| Orange     | 1     |
| Orange     | 23    |
| Peach      | 4     |
| Strawberry | 19    |

fruitsforsale

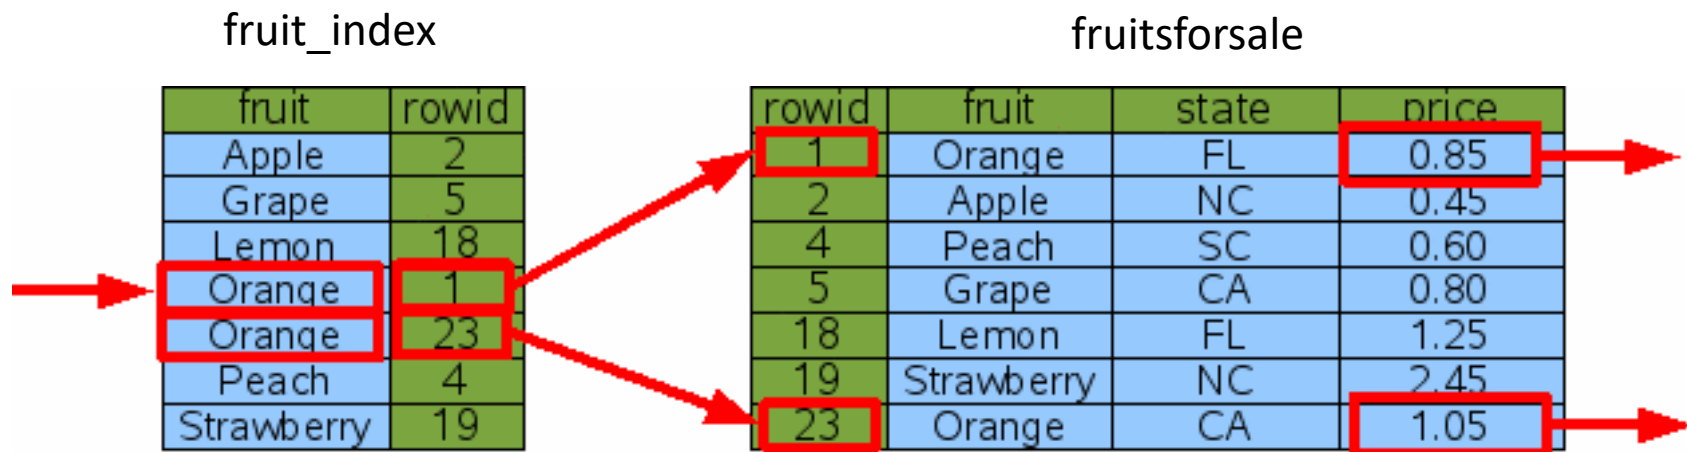
| rowid | fruit      | state | price |
|-------|------------|-------|-------|
| 1     | Orange     | FL    | 0.85  |
| 2     | Apple      | NC    | 0.45  |
| 4     | Peach      | SC    | 0.60  |
| 5     | Grape      | CA    | 0.80  |
| 18    | Lemon      | FL    | 1.25  |
| 19    | Strawberry | NC    | 2.45  |
| 23    | Orange     | CA    | 1.05  |

1. Do a binary search in fruit\_index for 'Peach', then we find the rowid.
2. Do a binary search in fruitsforsale for that rowid, then we can return the price.

# What is the Big O Notation for a Lookup with an Index?

- $n$
- $\log(n)$
- $\log(n) + \log(n)$
- Uhhhh...

```
SELECT price FROM fruitsforsale WHERE fruit = 'Orange';
```



1. Do a binary search in **fruit\_index** for 'Orange', then we find the rowid.
2. Do a binary search in **fruitsforsale** for that rowid, then we can return the price.
3. Go back to **fruit\_index** and check if the next row is also an orange, if so look up its row in **fruitsforsale**.
4. Repeat until there are no more 'Orange's.

```
SELECT price FROM fruitsforsale WHERE fruit = 'Orange' AND
state = 'CA';
```

fruit\_index

| fruit      | rowid |
|------------|-------|
| Apple      | 2     |
| Grape      | 5     |
| Lemon      | 18    |
| Orange     | 1     |
| Orange     | 23    |
| Peach      | 4     |
| Strawberry | 19    |

fruitsforsale

| rowid | fruit      | state | price |
|-------|------------|-------|-------|
| 1     | Orange     | FL    | 0.85  |
| 2     | Apple      | NC    | 0.45  |
| 4     | Peach      | SC    | 0.60  |
| 5     | Grape      | CA    | 0.80  |
| 18    | Lemon      | FL    | 1.25  |
| 19    | Strawberry | NC    | 2.45  |
| 23    | Orange     | CA    | 1.05  |

Same process as before, but we have to exclude some of the 'Orange' rows if the state doesn't match 'CA'.

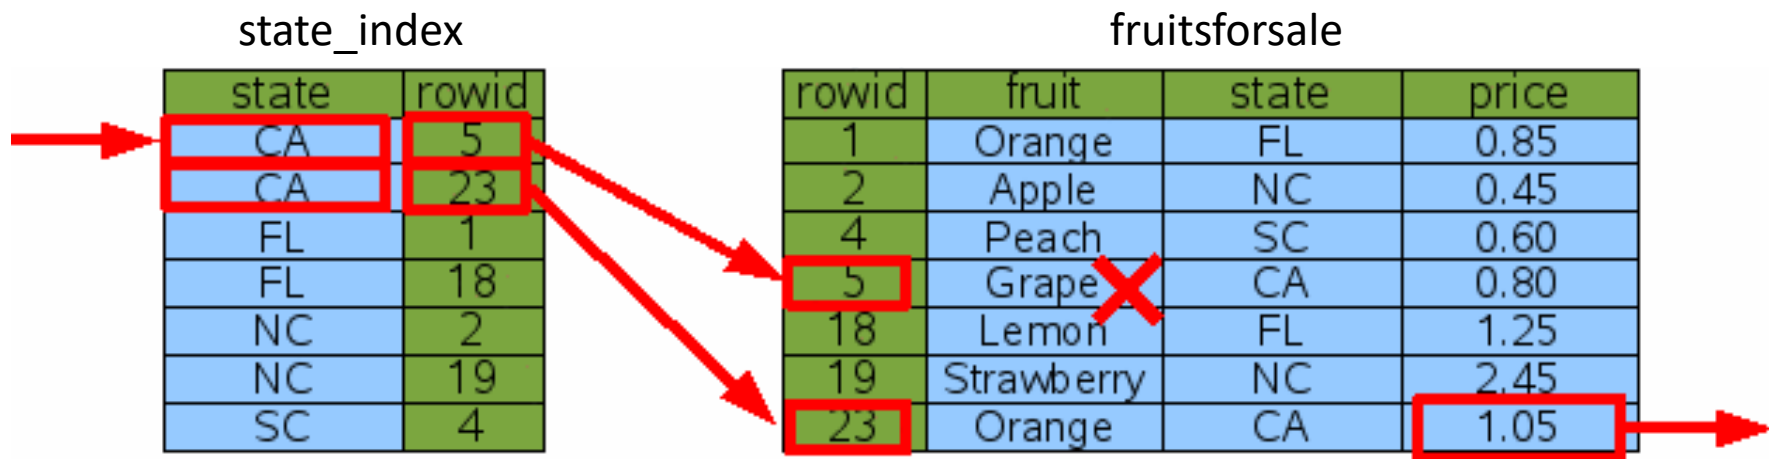


```
CREATE INDEX state_index ON fruitsforsale(state);
```

state\_index

| state | rowid |
|-------|-------|
| CA    | 5     |
| CA    | 23    |
| FL    | 1     |
| FL    | 18    |
| NC    | 2     |
| NC    | 19    |
| SC    | 4     |

```
SELECT price FROM fruitsforsale WHERE fruit =
'Orange' AND state = 'CA';
```



Using the state\_index instead of the fruit\_index follows the same process.

# If you have Multiple Indices, Which one is Faster?

- The index with the most rows
- The index with the least duplicates
- The index that is sorted
- Depends

```
CREATE INDEX fruit_state_index ON fruitsforsale(fruit,  
state);
```

fruit\_state\_index

| fruit      | state | rowid |
|------------|-------|-------|
| Apple      | NC    | 2     |
| Grape      | CA    | 5     |
| Lemon      | FL    | 18    |
| Orange     | CA    | 23    |
| Orange     | FL    | 1     |
| Peach      | SC    | 4     |
| Strawberry | NC    | 19    |

Multi-column index that is sorted according to the first column (ties broken by subsequent columns).

```
SELECT price FROM fruitsforsale WHERE  
fruit = 'Orange' AND state = 'CA';
```

fruit\_state\_index

| fruit      | state | rowid |
|------------|-------|-------|
| Apple      | NC    | 2     |
| Grape      | CA    | 5     |
| Lemon      | FL    | 18    |
| Orange     | CA    | 23    |
| Orange     | FL    | 1     |
| Peach      | SC    | 4     |
| Strawberry | NC    | 19    |

fruitsforsale

| rowid | fruit      | state | price |
|-------|------------|-------|-------|
| 1     | Orange     | FL    | 0.85  |
| 2     | Apple      | NC    | 0.45  |
| 4     | Peach      | SC    | 0.60  |
| 5     | Grape      | CA    | 0.80  |
| 18    | Lemon      | FL    | 1.25  |
| 19    | Strawberry | NC    | 2.45  |
| 23    | Orange     | CA    | 1.05  |

Using the fruit\_state\_index allows only finding the rows we want.

```
SELECT price FROM fruitsforsale WHERE fruit
= 'Peach';
```

fruit\_state\_index

| fruit      | state | rowid |
|------------|-------|-------|
| Apple      | NC    | 2     |
| Grape      | CA    | 5     |
| Lemon      | FL    | 18    |
| Orange     | CA    | 23    |
| Orange     | FL    | 1     |
| Peach      | SC    | 4     |
| Strawberry | NC    | 19    |

fruitsforsale

| rowid | fruit      | state | price |
|-------|------------|-------|-------|
| 1     | Orange     | FL    | 0.85  |
| 2     | Apple      | NC    | 0.45  |
| 4     | Peach      | SC    | 0.60  |
| 5     | Grape      | CA    | 0.80  |
| 18    | Lemon      | FL    | 1.25  |
| 19    | Strawberry | NC    | 2.45  |
| 23    | Orange     | CA    | 1.05  |

fruit\_state\_index has all of utility fruit\_index had, we can just ignore the state if it isn't needed.

```
CREATE INDEX fruit_state_price_index ON  
fruitsforsale(fruit, state, price);
```

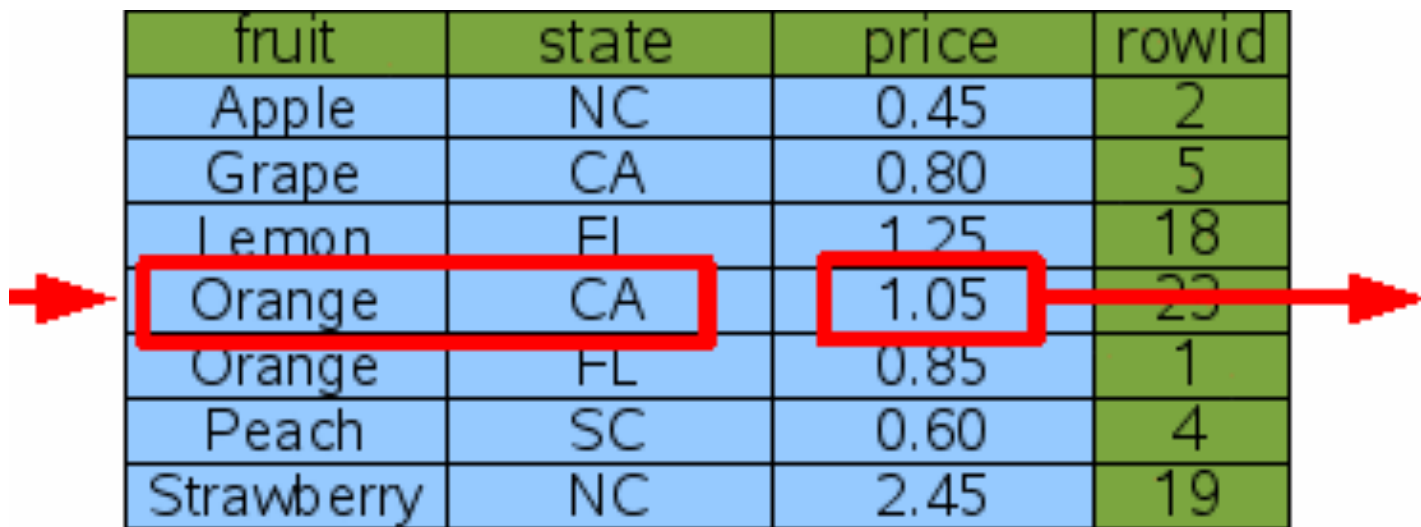
fruit\_state\_price\_index

| fruit      | state | price | rowid |
|------------|-------|-------|-------|
| Apple      | NC    | 0.45  | 2     |
| Grape      | CA    | 0.80  | 5     |
| Lemon      | FL    | 1.25  | 18    |
| Orange     | CA    | 1.05  | 23    |
| Orange     | FL    | 0.85  | 1     |
| Peach      | SC    | 0.60  | 4     |
| Strawberry | NC    | 2.45  | 19    |

This is called a covering index - it has all of the columns used in the SELECT statement, including the output ('price').

```
SELECT price FROM fruitsforsale WHERE  
fruit='Orange' AND states = 'CA';
```

fruit\_state\_price\_index



The diagram shows a table with 4 columns: fruit, state, price, and rowid. The first row is highlighted in green. The second row is highlighted in blue. The third row is highlighted in green. The fourth row is highlighted in blue. The fifth row is highlighted in green. The sixth row is highlighted in blue. The seventh row is highlighted in green. The eighth row is highlighted in blue. The ninth row is highlighted in green. The tenth row is highlighted in blue. The eleventh row is highlighted in green. The twelfth row is highlighted in blue. The thirteenth row is highlighted in green. The fourteenth row is highlighted in blue. The fifteenth row is highlighted in green. The sixteenth row is highlighted in blue. The seventeenth row is highlighted in green. The eighteenth row is highlighted in blue. The nineteenth row is highlighted in green. The twentieth row is highlighted in blue. The twenty-first row is highlighted in green. The twenty-second row is highlighted in blue. The twenty-third row is highlighted in green. The twenty-fourth row is highlighted in blue. The twenty-fifth row is highlighted in green. The twenty-sixth row is highlighted in blue. The twenty-seventh row is highlighted in green. The twenty-eighth row is highlighted in blue. The twenty-ninth row is highlighted in green. The thirtieth row is highlighted in blue. The thirty-first row is highlighted in green. The thirty-second row is highlighted in blue. The thirty-third row is highlighted in green. The thirty-fourth row is highlighted in blue. The thirty-fifth row is highlighted in green. The thirty-sixth row is highlighted in blue. The thirty-seventh row is highlighted in green. The thirty-eighth row is highlighted in blue. The thirty-ninth row is highlighted in green. The fortieth row is highlighted in blue. The forty-first row is highlighted in green. The forty-second row is highlighted in blue. The forty-third row is highlighted in green. The forty-fourth row is highlighted in blue. The forty-fifth row is highlighted in green. The forty-sixth row is highlighted in blue. The forty-seventh row is highlighted in green. The forty-eighth row is highlighted in blue. The forty-ninth row is highlighted in green. The fiftieth row is highlighted in blue. The fifty-first row is highlighted in green. The fifty-second row is highlighted in blue. The fifty-third row is highlighted in green. The fifty-fourth row is highlighted in blue. The fifty-fifth row is highlighted in green. The fifty-sixth row is highlighted in blue. The fifty-seventh row is highlighted in green. The fifty-eighth row is highlighted in blue. The fifty-ninth row is highlighted in green. The sixtieth row is highlighted in blue. The sixty-first row is highlighted in green. The sixty-second row is highlighted in blue. The sixty-third row is highlighted in green. The sixty-fourth row is highlighted in blue. The sixty-fifth row is highlighted in green. The sixty-sixth row is highlighted in blue. The sixty-seventh row is highlighted in green. The sixty-eighth row is highlighted in blue. The sixty-ninth row is highlighted in green. The seventieth row is highlighted in blue. The seventy-first row is highlighted in green. The seventy-second row is highlighted in blue. The seventy-third row is highlighted in green. The seventy-fourth row is highlighted in blue. The seventy-fifth row is highlighted in green. The seventy-sixth row is highlighted in blue. The seventy-seventh row is highlighted in green. The seventy-eighth row is highlighted in blue. The seventy-ninth row is highlighted in green. The eightieth row is highlighted in blue. The eighty-first row is highlighted in green. The eighty-second row is highlighted in blue. The eighty-third row is highlighted in green. The eighty-fourth row is highlighted in blue. The eighty-fifth row is highlighted in green. The eighty-sixth row is highlighted in blue. The eighty-seventh row is highlighted in green. The eighty-eighth row is highlighted in blue. The eighty-ninth row is highlighted in green. The ninetieth row is highlighted in blue. The ninety-first row is highlighted in green. The ninety-second row is highlighted in blue. The ninety-third row is highlighted in green. The ninety-fourth row is highlighted in blue. The ninety-fifth row is highlighted in green. The ninety-sixth row is highlighted in blue. The ninety-seventh row is highlighted in green. The ninety-eighth row is highlighted in blue. The ninety-ninth row is highlighted in green. The hundredth row is highlighted in blue. The hundred and first row is highlighted in green. The hundred and second row is highlighted in blue. The hundred and third row is highlighted in green. The hundred and fourth row is highlighted in blue. The hundred and fifth row is highlighted in green. The hundred and sixth row is highlighted in blue. The hundred and seventh row is highlighted in green. The hundred and eighth row is highlighted in blue. The hundred and ninth row is highlighted in green. The hundred and tenth row is highlighted in blue. The hundred and eleventh row is highlighted in green. The hundred and twelfth row is highlighted in blue. The hundred and thirteenth row is highlighted in green. The hundred and fourteenth row is highlighted in blue. The hundred and fifteenth row is highlighted in green. The hundred and sixteenth row is highlighted in blue. The hundred and seventeenth row is highlighted in green. The hundred and eighteenth row is highlighted in blue. The hundred and nineteenth row is highlighted in green. The hundred and twentieth row is highlighted in blue. The hundred and twenty-first row is highlighted in green. The hundred and twenty-second row is highlighted in blue. The hundred and twenty-third row is highlighted in green. The hundred and twenty-fourth row is highlighted in blue. The hundred and twenty-fifth row is highlighted in green. The hundred and twenty-sixth row is highlighted in blue. The hundred and twenty-seventh row is highlighted in green. The hundred and twenty-eighth row is highlighted in blue. The hundred and twenty-ninth row is highlighted in green. The hundred and thirtieth row is highlighted in blue. The hundred and thirty-first row is highlighted in green. The hundred and thirty-second row is highlighted in blue. The hundred and thirty-third row is highlighted in green. The hundred and thirty-fourth row is highlighted in blue. The hundred and thirty-fifth row is highlighted in green. The hundred and thirty-sixth row is highlighted in blue. The hundred and thirty-seventh row is highlighted in green. The hundred and thirty-eighth row is highlighted in blue. The hundred and thirty-ninth row is highlighted in green. The hundred and fortieth row is highlighted in blue. The hundred and forty-first row is highlighted in green. The hundred and forty-second row is highlighted in blue. The hundred and forty-third row is highlighted in green. The hundred and forty-fourth row is highlighted in blue. The hundred and forty-fifth row is highlighted in green. The hundred and forty-sixth row is highlighted in blue. The hundred and forty-seventh row is highlighted in green. The hundred and forty-eighth row is highlighted in blue. The hundred and forty-ninth row is highlighted in green. The hundred and fiftieth row is highlighted in blue. The hundred and fifty-first row is highlighted in green. The hundred and fifty-second row is highlighted in blue. The hundred and fifty-third row is highlighted in green. The hundred and fifty-fourth row is highlighted in blue. The hundred and fifty-fifth row is highlighted in green. The hundred and fifty-sixth row is highlighted in blue. The hundred and fifty-seventh row is highlighted in green. The hundred and fifty-eighth row is highlighted in blue. The hundred and fifty-ninth row is highlighted in green. The hundred and sixtieth row is highlighted in blue. The hundred and sixty-first row is highlighted in green. The hundred and sixty-second row is highlighted in blue. The hundred and sixty-third row is highlighted in green. The hundred and sixty-fourth row is highlighted in blue. The hundred and sixty-fifth row is highlighted in green. The hundred and sixty-sixth row is highlighted in blue. The hundred and sixty-seventh row is highlighted in green. The hundred and sixty-eighth row is highlighted in blue. The hundred and sixty-ninth row is highlighted in green. The hundred and seventieth row is highlighted in blue. The hundred and seventy-first row is highlighted in green. The hundred and seventy-second row is highlighted in blue. The hundred and seventy-third row is highlighted in green. The hundred and seventy-fourth row is highlighted in blue. The hundred and seventy-fifth row is highlighted in green. The hundred and seventy-sixth row is highlighted in blue. The hundred and seventy-seventh row is highlighted in green. The hundred and seventy-eighth row is highlighted in blue. The hundred and seventy-ninth row is highlighted in green. The hundred and eightieth row is highlighted in blue. The hundred and eighty-first row is highlighted in green. The hundred and eighty-second row is highlighted in blue. The hundred and eighty-third row is highlighted in green. The hundred and eighty-fourth row is highlighted in blue. The hundred and eighty-fifth row is highlighted in green. The hundred and eighty-sixth row is highlighted in blue. The hundred and eighty-seventh row is highlighted in green. The hundred and eighty-eighth row is highlighted in blue. The hundred and eighty-ninth row is highlighted in green. The hundred and ninetieth row is highlighted in blue. The hundred and ninety-first row is highlighted in green. The hundred and ninety-second row is highlighted in blue. The hundred and ninety-third row is highlighted in green. The hundred and ninety-fourth row is highlighted in blue. The hundred and ninety-fifth row is highlighted in green. The hundred and ninety-sixth row is highlighted in blue. The hundred and ninety-seventh row is highlighted in green. The hundred and ninety-eighth row is highlighted in blue. The hundred and ninety-ninth row is highlighted in green. The thousandth row is highlighted in blue.

| fruit      | state | price | rowid |
|------------|-------|-------|-------|
| Apple      | NC    | 0.45  | 2     |
| Grape      | CA    | 0.80  | 5     |
| Lemon      | FL    | 1.25  | 18    |
| Orange     | CA    | 1.05  | 23    |
| Orange     | FL    | 0.85  | 1     |
| Peach      | SC    | 0.60  | 4     |
| Strawberry | NC    | 2.45  | 19    |

A covering index for a query doesn't have to consult the original table, because all of the information is in the index.

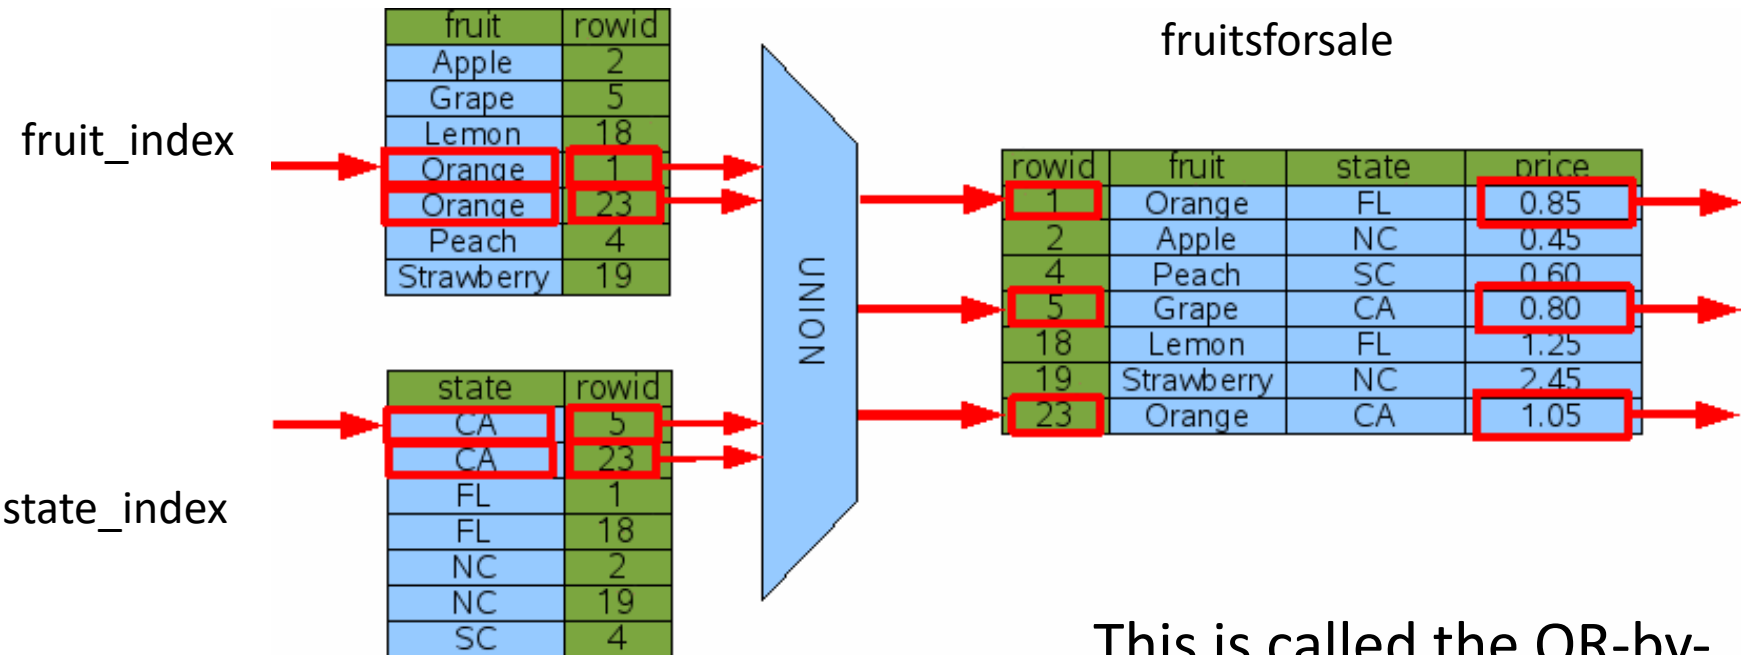
This means the second binary lookup isn't done.



# What is the Big O Notation for a Lookup with a Covering Index?

- $n$
- $\log(n)$
- $\log(n) + \log(n)$
- Uhhhh...

SELECT price FROM fruitsforsale WHERE  
 fruit='Orange' OR state='CA';



This is called the OR-by-  
 UNION technique.

```
SELECT * FROM fruitsforsale ORDER BY fruit;
```

fruitsforsale



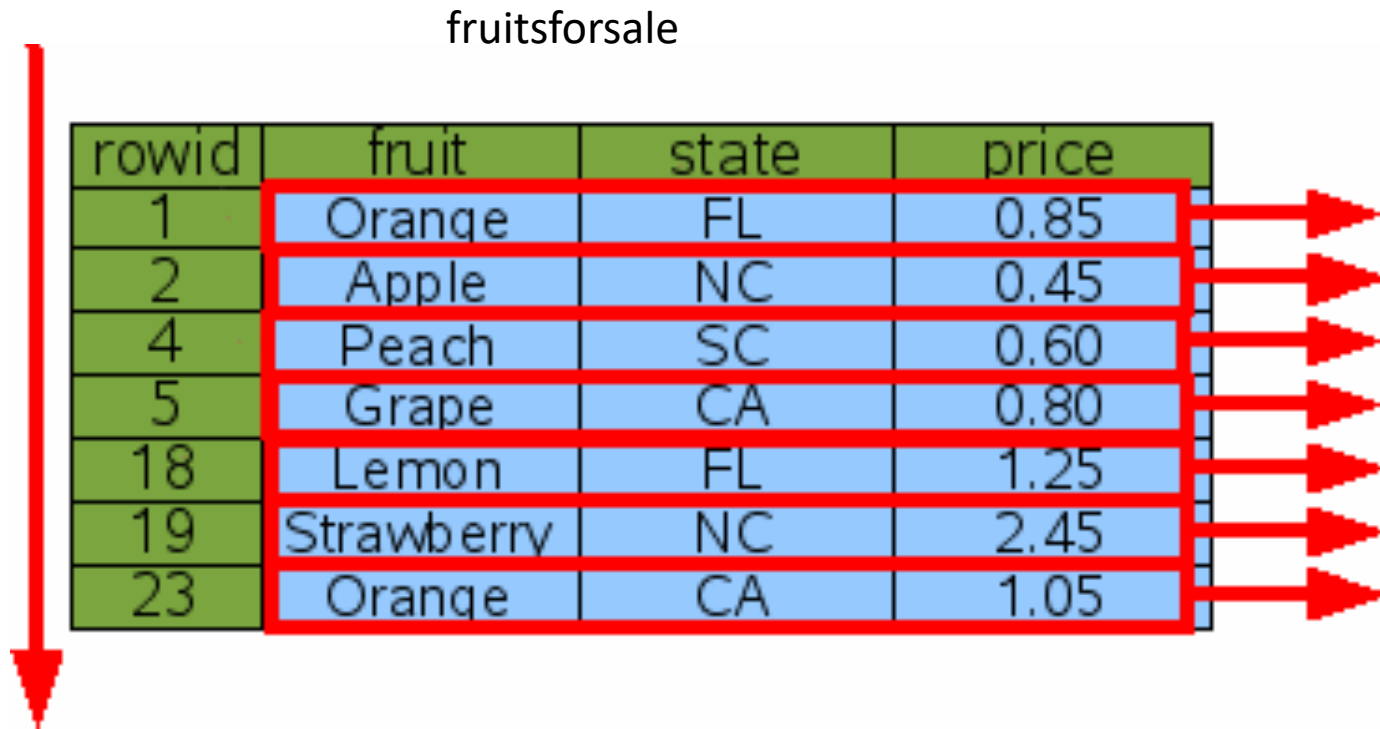
If we didn't have an index, every row is passed to a sorter function.

# What is the Big O Notation for Sorting Without an Index?

- $n$
- $\log(n)$
- $n \log(n)$
- Uhhhh...

```
SELECT * FROM fruitsforsale ORDER BY rowid;
```

fruitsforsale



| rowid | fruit      | state | price |
|-------|------------|-------|-------|
| 1     | Orange     | FL    | 0.85  |
| 2     | Apple      | NC    | 0.45  |
| 4     | Peach      | SC    | 0.60  |
| 5     | Grape      | CA    | 0.80  |
| 18    | Lemon      | FL    | 1.25  |
| 19    | Strawberry | NC    | 2.45  |
| 23    | Orange     | CA    | 1.05  |

No sorting needed, just return the rows as they are.

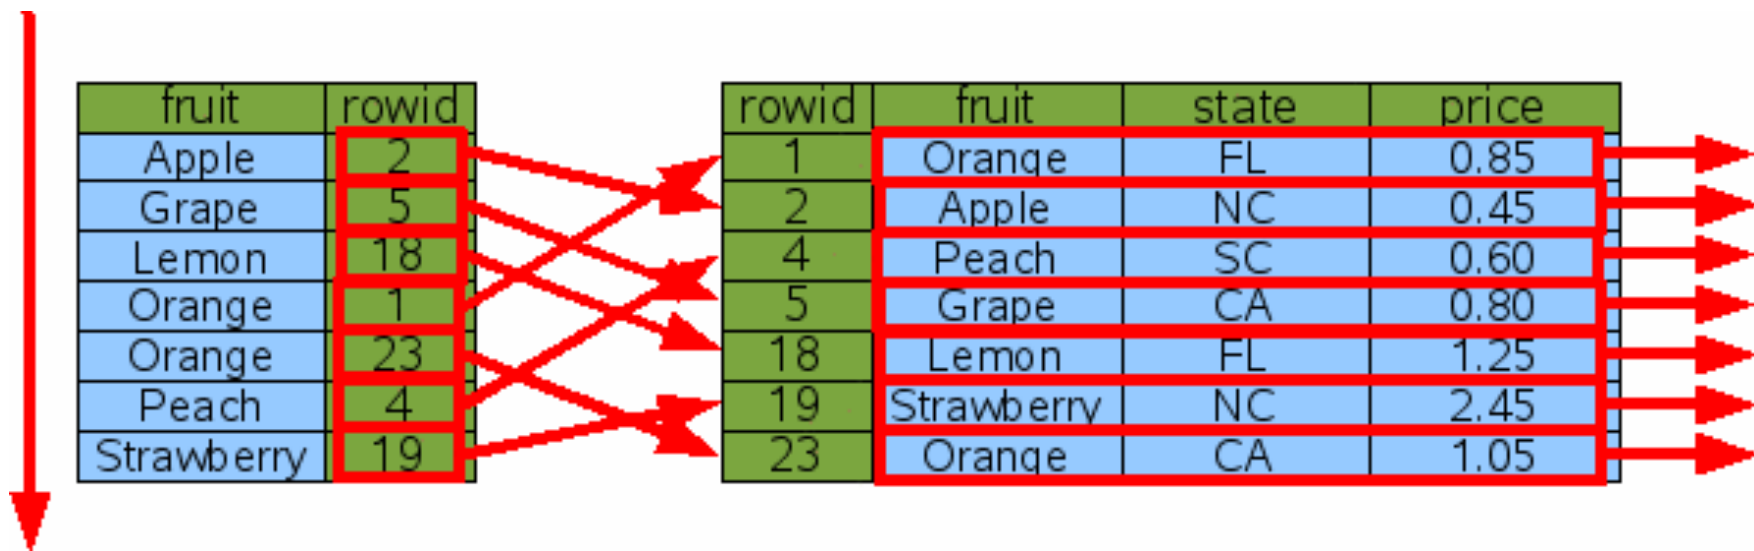
Can you imagine what happens with:

```
"SELECT * FROM fruitsforsale ORDER BY rowid DESC;"
```

SELECT \* FROM fruitsforsale ORDER BY fruit;

fruits\_index

fruitsforsale



Using the index, we don't have to pass everything to a sorted function

# What is the Big O Notation for Sorting with an Index?

- $n$
- $\log(n)$
- $n \log(n)$
- Uhhhh...

```
SELECT price FROM fruitsforsale WHERE  
fruit='Orange' ORDER BY state;
```

fruits\_state\_index

| fruit      | state | rowid |
|------------|-------|-------|
| Apple      | NC    | 2     |
| Grape      | CA    | 5     |
| Lemon      | FL    | 18    |
| Orange     | CA    | 23    |
| Orange     | FL    | 1     |
| Peach      | SC    | 4     |
| Strawberry | NC    | 19    |

fruitsforsale

| rowid | fruit      | state | price |
|-------|------------|-------|-------|
| 1     | Orange     | FL    | 0.85  |
| 2     | Apple      | NC    | 0.45  |
| 4     | Peach      | SC    | 0.60  |
| 5     | Grape      | CA    | 0.80  |
| 18    | Lemon      | FL    | 1.25  |
| 19    | Strawberry | NC    | 2.45  |
| 23    | Orange     | CA    | 1.05  |

- If we are sorting and filtering, using a appropriate index is very fast.
1. Find rows with fruit='Orange' (left most column in index)
  2. Output the price for each record (they are already sorted by state).



```
SELECT * FROM fruitsforsale WHERE  
fruit='Orange' ORDER BY state;
```

fruits\_state\_price\_index

| fruit      | state | price | rowid |
|------------|-------|-------|-------|
| Apple      | NC    | 0.45  | 2     |
| Grape      | CA    | 0.80  | 5     |
| Lemon      | FL    | 1.25  | 18    |
| Orange     | CA    | 1.05  | 23    |
| Orange     | FL    | 0.85  | 1     |
| Peach      | SC    | 0.60  | 4     |
| Strawberry | NC    | 2.45  | 19    |

Using the appropriate covering index, we don't even have to do a lookup on the underlying table.

# That's it for today

- Questions?